

Machine-Level Programming I: Basics

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



1

1

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



2

2

Intel x86 Processors

- Dominate laptop/desktop/server market
- Evolutionary design
 - Backwards compatible up until 8086, introduced in 1978
 - Added more features as time goes on
- Complex instruction set computer (CISC)
 - Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
 - Hard to match performance of Reduced Instruction Set Computers (RISC)
 - But, Intel has done just that!
 - In terms of speed. Less so for low power.

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



3

3

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



4

4

Definitions

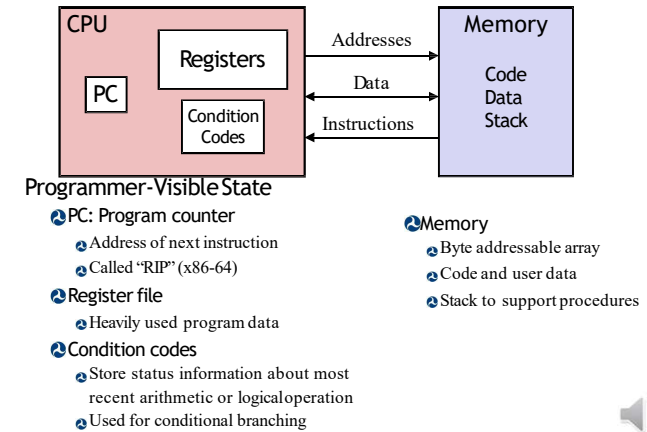
- **Instruction Set Architecture (ISA):** The parts of a processor design that one needs to understand or write assembly/machine code.
 - Examples: instruction set specification, registers.
- **Microarchitecture:** Implementation of the architecture.
 - Examples: cache sizes and core frequency.
- **Code Forms:**
 - **Machine Code:** The byte-level programs that a processor executes
 - **Assembly Code:** A text representation of machine code
- **Example ISAs:**
 - Intel: x86, IA32, Itanium, x86-64
 - ARM: Used in almost all mobile phones

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



5

Assembly/Machine Code View



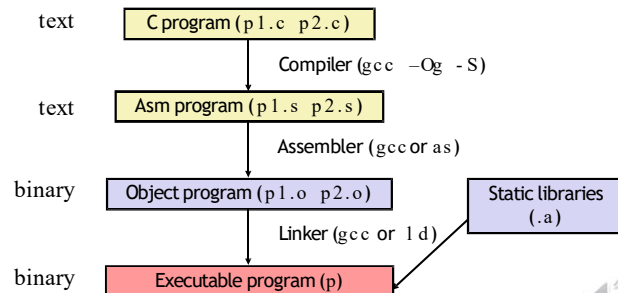
Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



6

Turning C into Object Code

- Code in files `p1.c p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
 - Use basic optimizations (-Og) [New to recent versions of GCC]
 - Put resulting binary in file



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



7

Assembly Characteristics: Data Types

- "Integer" data of 1, 2, 4, or 8 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- Code: Byte sequences encoding series of instructions
- No aggregate types such as arrays or structures
 - Just contiguously allocated bytes in memory

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



8

Assembly Characteristics: Operations

- ⚙ Perform arithmetic function on register or memory data
- ⚙ Transfer data between memory and register
 - ⚙ Load data from memory into register
 - ⚙ Store register data into memory
- ⚙ Transfer control
 - ⚙ Unconditional jumps to/from procedures
 - ⚙ Conditional branches

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



9

Object Code

Code for `sumstore`

```
0x0400595:
0x53
0x48
0x89
0xd3
0xe8
0xf2
0xff
0xff
0xff
0x48
0x89
0x03
0x5b
0xc3
```

- Total of 14 bytes
- Each instruction 1, 3, or 5 bytes
- Starts at address 0x0400595

⚙ Assembler

- ⚙ Translates `.S` into `.o`
- ⚙ Binary encoding of each instruction
- ⚙ Nearly-complete image of executable code
- ⚙ Missing linkages between code in different files

⚙ Linker

- ⚙ Resolves references between files
- ⚙ Combines with static run-time libraries
 - ⚙ E.g., code for `malloc`, `printf`
- ⚙ Some libraries are dynamically linked
 - ⚙ Linking occurs when program begins execution

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



10

Machine Instruction Example

```
*dest = t;
```

⚙ C Code

- ⚙ Store value `t` where designated by `dest`

```
movq %rax, (%rbx)
```

⚙ Assembly

- ⚙ Move 8-byte value to memory
 - ⚙ Quad words in x86-64 parlance
- ⚙ Operands:
 - `t`: Register `%rax`
 - `dest`: Register `%rbx`
 - `*dest`: Memory `M[%rbx]`

```
0x40059e: 48 89 03
```

⚙ Object Code

- ⚙ 3-byte instruction
- ⚙ Stored at address `0x40059e`

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



11

Disassembling Object Code

Disassembled

```
000000000400595 <sumstore>:
400595: 53          push    %rbx
400596: 48 89 d3    mov     %rdx,%rbx
400599: e8 f2 ff ff callq   400590 <plus>
40059e: 48 89 03    mov     %rax,(%rbx)
4005a1: 5b         pop     %rbx
4005a2: c3         retq
```

⚙ Disassembler

```
objdump -d sum
```

- ⚙ Useful tool for examining object code
- ⚙ Analyzes bit pattern of series of instructions
- ⚙ Produces approximate rendition of assembly code
- ⚙ Can be run on either `a.out` (complete executable) or `.o` file

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



12

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move**
- Arithmetic & logical operations

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



13

13

x86-64 Integer Registers

%rax	%eax	%r8	%r8d
%rbx	%ebx	%r9	%r9d
%rcx	%ecx	%r10	%r10d
%rdx	%edx	%r11	%r11d
%rsi	%esi	%r12	%r12d
%rdi	%edi	%r13	%r13d
%rsp	%esp	%r14	%r14d
%rbp	%ebp	%r15	%r15d

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



14

14

Moving Data

Moving Data

movq Source, Dest;

Operand Types

- Immediate:** Constant integer data

- Example: \$0x400, \$-533
- Like C constant, but prefixed with '\$'
- Encoded with 1, 2, or 4 bytes

- Register:** One of 16 integer registers

- Example: %rax, %r13
- But %rsp reserved for special use
- Others have special uses for particular instructions

- Memory:** 8 consecutive bytes of memory at address given by register

- Simplest example: (%rax)
- Various other "address modes"

%rax
%rcx
%rdx
%rbx
%rsi
%rdi
%rsp
%rbp
%rN

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



15

15

movq Operand Combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x4, %rax	temp = 0x4;
		Mem	movq \$-147, (%rax)	*p = -147;
	Reg	Reg	movq %rax, %rdx	temp2 = temp1;
		Mem	movq %rax, (%rdx)	*p = temp;
	Mem	Reg	movq (%rax), %rdx	temp = *p;
		Mem		

Cannot do memory-memory transfer with a single instruction

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



16

16


Understanding Swap()

```
void swap
(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Registers

%rdi	
%rsi	
%rax	
%rdx	

Memory



Register	Value
%rdi	xp
%rsi	yp
%rax	t0
%rdx	t1

swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)    # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



17

17

Simple Memory Addressing Modes

Normal (R) Mem[Reg[R]]

- Register R specifies memory address
- Aha! Pointer dereferencing in C

```
movq    (%rcx), %rax
```

Displacement D(R) Mem[Reg[R]+D]

- Register R specifies start of memory region
- Constant displacement D specifies offset

```
movq    8(%rbp), %rdx
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



18

18

Complete Memory Addressing Modes

Most General Form

$D(Rb, Ri, S)$ Mem[Reg[Rb]+S*Reg[Ri]+D]

- D: Constant "displacement" 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for %rsp
- S: Scale: 1, 2, 4, or 8 (why these numbers?)

Special Cases

(Rb, Ri) Mem[Reg[Rb]+Reg[Ri]]
 D(Rb, Ri) Mem[Reg[Rb]+Reg[Ri]+D]
 (Rb, Ri, S) Mem[Reg[Rb]+S*Reg[Ri]]

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



19

19

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



20

20

Address Computation Instruction

- `leaq Src, Dst`
 - Src is address mode expression
 - Set Dst to address denoted by expression
- Uses
 - Computing addresses without a memory reference
 - E.g., translation of `p = &x[i];`
 - Computing arithmetic expressions of the form `x+k*y`
 - $k=1, 2, 4, \text{ or } 8$

Example

```
long ml2(long x)
{
    return x*12;
}
```

Converted to ASM by compiler:

```
leaq (%rdi,%rdi,2), %rax # t <- x+x*2
salq $2, %rax           # return t<<2
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

21

21

Some Arithmetic Operations

Two Operand Instructions:

Format	Computation	
<code>addq Src, Dest</code>	<code>Dest = Dest + Src</code>	
<code>subq Src, Dest</code>	<code>Dest = Dest - Src</code>	
<code>imulq Src, Dest</code>	<code>Dest = Dest * Src</code>	
<code>salq Src, Dest</code>	<code>Dest = Dest << Src</code>	Also called <code>shlq</code> Arithmetic Logical
<code>sarq Src, Dest</code>	<code>Dest = Dest >> Src</code>	
<code>shrq Src, Dest</code>	<code>Dest = Dest >> Src</code>	
<code>xorq Src, Dest</code>	<code>Dest = Dest ^ Src</code>	
<code>and Src, Dest</code>	<code>Dest = Dest & Src</code>	
<code>q Src, Dest</code>	<code>Dest = Dest Src</code>	

- Watch out for argument order!
- No distinction between signed and unsigned int (why?)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

22

22

Some Arithmetic Operations

- One Operand Instructions

<code>incq Dest</code>	<code>Dest = Dest + 1</code>
<code>decq Dest</code>	<code>Dest = Dest - 1</code>
<code>negq Dest</code>	<code>Dest = -Dest</code>
<code>notq Dest</code>	<code>Dest = ~Dest</code>
- See book for more instructions

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

23

23

Understanding Arithmetic Expression Example

```
long arith
(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

```
arith:
    leaq    (%rdi,%rsi), %rax    # t1
    addq    %rdx, %rax          # t2
    leaq    (%rsi,%rsi,2), %rdx
    salq    $4, %rdx            # t4
    leaq    4(%rdi,%rdx), %rcx   # t5
    imulq    %rcx, %rax          # rval
    ret
```

<code>%rdi</code>	Argument x
<code>%rsi</code>	Argument y
<code>%rdx</code>	Argument z
<code>%rax</code>	t1, t2, rval
<code>%rdx</code>	t4
<code>%rcx</code>	t5

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

24

24

Machine Programming I: Summary

- History of Intel processors and architectures
 - Evolutionary design leads to many quirks and artifacts
- C, assembly, machine code
 - New forms of visible state: program counter, registers, ...
 - Compiler must transform statements, expressions, procedures into low-level instruction sequences
- Assembly Basics: Registers, operands, move
 - The x86-64 move instructions cover wide range of data movement forms
- Arithmetic
 - Compiler will figure out different instruction combinations to carry out computation

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



25

25

Machine-Level Programming II: Control

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



26

26

Today

- Control: Condition codes
- Conditional branches
- Loops
- Switch Statements

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



27

27

Processor State (x86-64, Partial)

- Information about currently executing program

- Temporary data (%rax, ...)
- Location of runtime stack (%rsp)
- Location of current code control point (%rip, ...)
- Status of recent tests (CF, ZF, SF, OF)

Registers

%rax	%r8
%rbx	%r9
%rcx	%r10
%rdx	%r11
%rsi	%r12
%rdi	%r13
%rsp	%r14
%rbp	%r15

%rip Instruction pointer

CF ZF SF OF

Condition codes

Current stack top

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



28

28

Condition Codes (Implicit Setting)

Single bit registers

- CF Carry Flag (for unsigned) SF Sign Flag (for signed)
- ZF Zero Flag OF Overflow Flag (for signed)

Implicitly set (think of it as side effect) by arithmetic operations

Example: `addq Src, Dest` \leftrightarrow `t = a + b`

CF set if carry out from most significant bit (unsigned overflow)

ZF set if `t == 0`

SF set if `t < 0` (signed)

OF set if two's-complement (signed) overflow
`(a > 0 && b > 0 && t < 0) || (a < 0 && b < 0 && t > 0)`

Not set by `leaq` instruction

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



29

Condition Codes (Explicit Setting: Compare)

Explicit Setting by Compare Instruction

`cmpq Src2, Src1`

`cmpq b, a` like computing `a - b` without setting destination

CF set if carry out from most significant bit (used for unsigned comparisons)

ZF set if `a == b`

SF set if `(a - b) < 0` (signed)

OF set if two's-complement (signed) overflow
`(a > 0 && b < 0 && (a - b) < 0) || (a < 0 && b > 0 && (a - b) > 0)`

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



30

Condition Codes (Explicit Setting: Test)

Explicit Setting by Test instruction

`testq Src2, Src1`

`testq b, a` like computing `a & b` without setting destination

Sets condition codes based on value of `Src1` & `Src2`

Useful to have one of the operands be a mask

ZF set when `a & b == 0`

SF set when `a & b < 0`

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



31

Reading Condition Codes

SetX Instructions

Set low-order byte of destination to 0 or 1 based on combinations of condition codes

Does not alter remaining 7 bytes

SetX	Condition	Description
<code>sete</code>	ZF	Equal / Zero
<code>setne</code>	\sim ZF	Not Equal / Not Zero
<code>sets</code>	SF	Negative
<code>setns</code>	\sim SF	Nonnegative
<code>setg</code>	\sim (SF ^ OF) & \sim ZF	Greater (Signed)
<code>setge</code>	\sim (SF ^ OF)	Greater or Equal (Signed)
<code>setl</code>	(SF ^ OF)	Less (Signed)
<code>setle</code>	(SF ^ OF) ZF	Less or Equal (Signed)
<code>seta</code>	\sim CF & \sim ZF	Above (unsigned)
<code>setb</code>	CF	Below (unsigned)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



32

x86-64 Integer Registers

%rax	%al	%r8	%r8b
%rbx	%bl	%r9	%r9b
%rcx	%cl	%r10	%r10b
%rdx	%dl	%r11	%r11b
%rsi	%sil	%r12	%r12b
%rdi	%dil	%r13	%r13b
%rsp	%spl	%r14	%r14b
%rbp	%bpl	%r15	%r15b

Can reference low-order byte

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



33

Today

- Control: Condition codes
- Conditional branches
- Loops
- Switch Statements

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



34

Jumping

jX Instructions

- Jump to different part of code depending on condition codes

jX	Condition	Description
jmp	1	Unconditional
j e	ZF	Equal / Zero
j ne	~ZF	Not Equal / Not Zero
j s	SF	Negative
j ns	~SF	Nonnegative
j g	~(SF^OF)&~ZF	Greater (Signed)
j ge	~(SF^OF)	Greater or Equal (Signed)
j l	(SF^OF)	Less (Signed)
j le	(SF^OF) ZF	Less or Equal (Signed)
j a	~CF&~ZF	Above (unsigned)
j b	CF	Below (unsigned)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



35

Conditional Branch Example (Old Style)

Generation

```
shark> gcc -Og -S -fno-if-conversion control.c
```

```
long absdiff
(long x, long y)
{
    long result;
    if (x >
        y)
        result = x - y;
    else
        result = y - x;
    return result;
}
```

```
absdiff:
    cmpq    %rsi, %rdi    # x:y
    jle     .L4
    movq    %rdi, %rax
    subq    %rsi, %rax
    ret
.L4:
    # x <= y
    movq    %rsi, %rax
    subq    %rdi, %rax
    ret
```

%rdi Argument x
%rsi Argument y
%rax Return value

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



36

Expressing with Goto Code

- ⌚ C allows goto statement
- ⌚ Jump to position designated by label

```
long absdiff
(long x, long y)
{
    long result;
    if (x > y)
        result = x-y;
    else
        result = y-x;
    return result;
}
```

```
long absdiff_j
(long x, long y)
{
    long result;
    int ntest = x <= y;
    if (ntest) goto Else;
    result = x-y;
    goto Done;
Else:
    result = y-x;
Done:
    return result;
}
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



37

General Conditional Expression Translation (Using Branches)

C Code

```
val = Test ? Then_Expr : Else_Expr;
```

```
val = x > y ? x - y : y - x;
```

Goto Version

```
ntest = !Test;
if (ntest) goto Else;
val = Then_Expr;
goto Done;
Else:
    val = Else_Expr;
Done:
```

- ⌚ Create separate code regions for then & else expressions
- ⌚ Execute appropriate one

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



38

Using Conditional Moves

- ⌚ Conditional Move Instructions
 - ⌚ Instruction supports: if(Test) Dest ← Src
 - ⌚ Supported in post-1995 x86 processors
 - ⌚ GCC tries to use them
 - ⌚ But, only when known to be safe
- ⌚ Why?
 - ⌚ Branches are very disruptive to instruction flow through pipelines
 - ⌚ Conditional moves do not require control transfer

C Code

```
val = Test
? Then_Expr
: Else_Expr;
```

Goto Version

```
result = Then_Expr;
eval = Else_Expr; ntest = !Test;
if (ntest) result = eval;
return result;
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



39

Conditional Move Example

```
long absdiff
(long x, long y)
{
    long result;
    if (x > y)
        result = x-y;
    else
        result = y-x;
    return result;
}
```

%rdi	Argument x
%rsi	Argument y
%rax	Return value

```
absdiff:
    movq    %rdi, %rax    # x
    subq    %rsi, %rax    # result = x-y
    movq    %rsi, %rdx
    subq    %rdi, %rdx    # eval = y-x
    cmpq    %rsi, %rdi    # x:y
    cmovle  %rdx, %rax    # if <=, result = eval
    ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



40

Bad Cases for Conditional Move

Expensive Computations

```
val = Test(x) ? Hard1(x) : Hard2(x);
```

- Both values get computed
- Only makes sense when computations are very simple

Risky Computations

```
val = p ? *p : 0;
```

- Both values get computed
- May have undesirable effects

Computations with side effects

```
val = x > 0 ? x*=7 : x+=3;
```

- Both values get computed
- Must be side-effect free

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



17

41

Today

- Control: Condition codes
- Conditional branches
- Loops
- Switch Statements

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



18

42

"Do-While" Loop Compilation

Goto Version

```
long pcount_goto
(unsigned long x) {
    long result = 0;
loop:
    result += x & 0x1;
    x >>= 1;
    if(x) goto loop;
    return result;
}
```

%rdi Argument x
%rax result

```
movl    $0, %eax    # result = 0
.L2:    # loop:
movq    %rdi, %rdx
andl    $1, %edx    # t = x & 0x1 #
addq    %dx, %rax    result += t #
shrq    %rdi        x >>= 1
jne     .L2          # if (x) goto loop
rep; ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



20

43

General "Do-While" Translation

C Code

```
do
    Body
while ( Test );
```

Goto Version

```
loop:
    Body
    if (Test)
        goto loop
```

- Body: {
 Statement₁;
 Statement₂;
 ...
 Statement_n;
 }

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



21

44

General “While” Translation #2

While version

```
while (Test)
  Body
```

- “Do-while” conversion
- Used with -O1

Do-While Version

```
if (!Test)
  goto done;
do
  Body
while (Test);
done:
```

Goto Version

```
if (!Test)
  goto done;
loop:
  Body
  if (Test) goto loop;
done:
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

24

45

“For” Loop \square While Loop

For Version

```
for (Init; Test; Update)
  Body
```

While Version

```
Init;
while (Test) {
  Body
  Update;
}
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

27

46

Today

- Control: Condition codes
- Conditional branches
- Loops
- Switch Statements

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

47

Switch Statement Example

```
long switch_eg
(long x, long y, long z)
{
  long w = 1;
  switch(x) {
    case 1:
      w = y*z;
      break;
    case 2:
      w = y/z;
      /* Fall Through */
    case 3:
      w += z;
      break;
    case 5:
    case 6:
      w -= z;
      break;
    default:
      w = 2;
  }
  return w;
}
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

31

- Multiple case labels

- Here: 5 & 6

- Fall through cases

- Here: 2

- Missing cases

- Here: 4

48

Jump Table Structure

Switch Form

```
switch(x) {
  case val_0:
    Block0
  case val_1:
    Block1
    . . .
  case val_n-1:
    Blockn-1
}
```

Jump Table

```
jtab: Targ0
      Targ1
      Targ2
      .
      .
      Targn-1
```

Jump Targets

```
Targ0: Code Block 0
Targ1: Code Block 1
Targ2: Code Block 2
      .
      .
      Targn-1: Code Block n-1
```

Translation (ExtendedC)

```
goto *JTab[x];
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

32

49

Assembly Setup Explanation

Table Structure

- Each target requires 8 bytes
- Base address at .L4

Jump table

```
.section .rodata
.align 8
.L4:
.quad .L8 # x = 0
.quad .L3 # x = 1
.quad .L5 # x = 2
.quad .L9 # x = 3
.quad .L8 # x = 4
.quad .L7 # x = 5
.quad .L7 # x = 6
```

Jumping

- Direct:** `jmp .L8`
- Jump target is denoted by label .L8
- Indirect:** `jmp *.L4(,%rdi,8)`
- Start of jump table: .L4
- Must scale by factor of 8 (addresses are 8 bytes)
- Fetch target from effective Address: $.L4 + x * 8$
- Only for $0 \leq x \leq 6$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

35

50

Jump Table

Jump table

```
.section .rodata
.align 8
.L4:
.quad .L8 # x = 0
.quad .L3 # x = 1
.quad .L5 # x = 2
.quad .L9 # x = 3
.quad .L8 # x = 4
.quad .L7 # x = 5
.quad .L7 # x = 6
```

```
switch(x) {
  case 1: // .L3
    w = y*z;
    break;
  case 2: // .L5
    w = y/z;
    /* Fall Through */
  case 3: // .L9
    w += z;
    break;
  case 5:
  case 6: // .L7
    w -= z;
    break;
  default: // .L8
    w = 2;
}
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

36

51

Code Blocks (x == 2, x == 3)

```
long w = 1;
switch(x) {
  case 2:
    w = y/z;
    /* Fall Through */
  case 3:
    w += z;
    break;
  . . .
}
```

```
.L5:                # Case 2
    movq    %rsi, %rax
    cqto
    idivq   %rcx     # y/z
    jmp     .L6      # goto merge
.L9:                # Case 3
    movl    $1, %eax # w = 1
.L6:                # merge:
    addq    %rcx, %rax # w += z
    ret
```

%rdi Argument x
%rsi Argument y
%rdx Argument z
%rax Return value

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

39

52

Summarizing

- C Control
 - if-then-else
 - do-while
 - while, for
 - switch
- Assembler Control
 - Conditional jump
 - Conditional move
 - Indirect jump (via jump tables)
 - Compiler generates code sequence to implement more complex control
- Standard Techniques
 - Loops converted to do-while or jump-to-middle form
 - Large switch statements use jump tables

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



41

53

Machine-Level Programming III: Procedures

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




1

54

Mechanisms in Procedures

- Passing control
 - To beginning of procedure code
 - Back to return point
- Passing data
 - Procedure arguments
 - Return value
- Memory management
 - Allocate during procedure execution
 - Deallocate upon return
- Mechanisms all implemented with machine instructions
- x86-64 implementation of a procedure uses only those mechanisms required



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



2

55

Today

- Procedures
 - Stack Structure
 - Calling Conventions
 - Passing control
 - Passing data
 - Managing local data
 - Illustration of Recursion

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

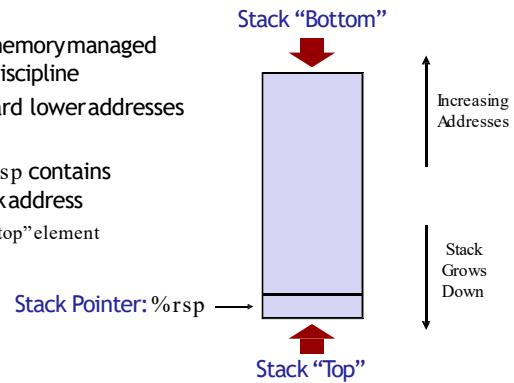


3

56

x86-64 Stack

- Region of memory managed with stack discipline
- Grows toward lower addresses
- Register `%rsp` contains lowest stack address
- address of "top" element



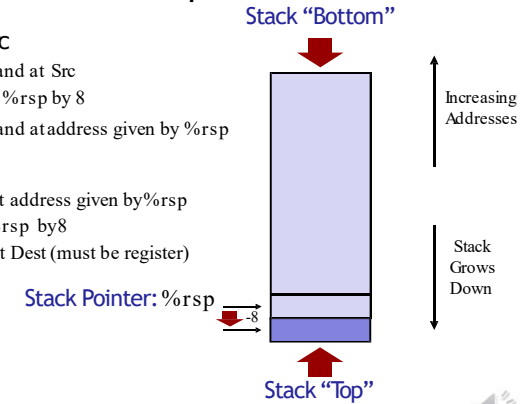
Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

4

57

x86-64 Stack: Push/Pop

- pushq Src**
 - Fetch operand at Src
 - Decrement `%rsp` by 8
 - Write operand at address given by `%rsp`
- popq Dest**
 - Read value at address given by `%rsp`
 - Increment `%rsp` by 8
 - Store value at Dest (must be register)



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

5

58

Code Examples

```
void multstore
(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

```
000000000400540 <multstore>:
400540: push    %rbx          # Save %rbx
400541: mov     %rdx,%rbx     # Save dest
400544: callq   400550 <mult2> # mult2(x,y)
400549: mov     %rax,(%rbx)    # Save at dest
40054c: pop     %rbx          # Restore %rbx
40054d: retq
```

```
long mult2
(long a, long b)
{
    long s = a * b;
    return s;
}
```

```
000000000400550 <mult2>:
400550: mov     %rdi,%rax     # a
400553: imul    %rsi,%rax     # a * b
400557: retq
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

8

59

Procedure Control Flow

- Use stack to support procedure call and return
- Procedure call: call label**
 - Push return address on stack
 - Jump to label
- Return address:**
 - Address of the next instruction right after call
 - Example from disassembly
- Procedure return: ret**
 - Pop address from stack
 - Jump to address

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

9

60

Procedure Data Flow

Registers

First 6 arguments

%rdi
%rsi
%rdx
%rcx
%r8
%r9

Return value

%rax

Stack

...
Argn
...
Arg8
Arg7

Only allocate stackspace when needed

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

15

61

Data Flow Examples

```
void multstore
(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

```
0000000000400540 <multstore>:
    # x in %rdi, y in %rsi, dest in %rdx
    ...
    400541: mov     %rdx,%rbx      # Save dest
    400544: callq  400550 <mult2>  # mult2(x, y)
    # t in %rax
    400549: mov     %rax, (%rbx)    # Save at dest
    ...
```

```
long mult2
(long a, long b)
{
    long s = a * b;
    return s;
}
```

```
0000000000400550 <mult2>:
    # a in %rdi, b in %rsi
    400550: mov     %rdi,%rax      # a
    400553: imul    %rsi,%rax      # a * b
    # s in %rax
    400557: retq                    # Return
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

16

62


Stack Frames

Contents

- Return information
- Local storage (if needed)
- Temporary space (if needed)

Frame Pointer: %rbp
(Optional)

Stack Pointer: %rsp



Stack "Top"

Management



- Space allocated when enter procedure
 - "Set-up" code
 - Includes push by call instruction
- Deallocated when return
 - "Finish" code
 - Includes pop by ret instruction

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

20

63

Example



%rbp

%rsp

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

25

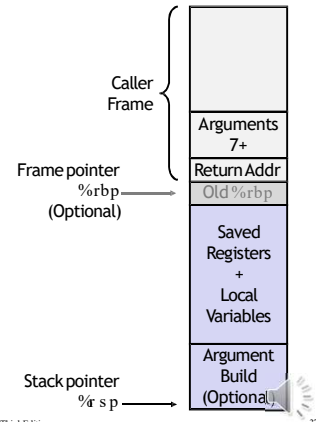
64

x86-64/Linux Stack Frame

- Current Stack Frame ("Top" to Bottom)
 - "Argument build:" Parameters for function about to call
 - Local variables If can't keep in registers
 - Saved register context
 - Old frame pointer (optional)

Caller Stack Frame

- Return address
 - Pushed by call instruction
- Arguments for this call




Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

32

65

x86-64 Linux Register Usage #1

- %rax
 - Return value
 - Also caller-saved
 - Can be modified by procedure
- %rdi, ..., %r9
 - Arguments
 - Also caller-saved
 - Can be modified by procedure
- %r10, %r11
 - Caller-saved
 - Can be modified by procedure



"Caller Saved"

Caller saves temporary values in its frame before the call




41

66

x86-64 Linux Register Usage #2

- %rbx, %r12, %r13, %r14
 - Callee-saved
 - Callee must save & restore
- %rbp
 - Callee-saved
 - Callee must save & restore
 - May be used as frame pointer
 - Can mix & match
- %rsp
 - Special form of callee save
 - Restored to original value upon exit from procedure



"Callee Saved"

Callee saves temporary values in its frame before using. Callee restores them before returning to caller




42

67

Callee-Saved Example #1


```
long call_incr2(long x) {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return x+v2;
}
```

Initial Stack Structure



```
call_incr2:
    pushq %rbx
    subq $16, %rsp
    movq %rdi, %rbx
    movq $15213, 8(%rsp)
    movl $3000, %esi
    leaq 8(%rsp), %rdi
    call incr
    addq %rbx, %rax
    addq $16, %rsp
    popq %rbx
    ret
```

Resulting Stack Structure



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



43


68

Recursive Function Register Save

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rdi	x	Argument



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

69

Observations About Recursion

Handled Without Special Consideration

- Stack frames mean that each function call has private storage
 - Saved registers & local variables
 - Saved return pointer
- Register saving conventions prevent one function call from corrupting another's data
 - Unless the C code explicitly does so (e.g., buffer overflow in Lecture 9)
- Stack discipline follows call / return pattern
 - If P calls Q, then Q returns before P
 - Last-In, First-Out

Also works for mutual recursion

- P calls Q; Q calls P

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

70

x86-64 Procedure Summary

Important Points


- Stack is the right data structure for procedure call / return
 - If P calls Q, then Q returns before P

Recursion (& mutual recursion) handled by normal calling conventions

- Can safely store values in local stack frame and in callee-saved registers
- Put function arguments at top of stack
- Result return in %rax

Pointers are addresses of values

- On stack or global



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


71

Machine-Level Programming IV: Data

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

72

Array Accessing Example



```
int get_digit
(zip_dig z, int digit)
{
    return z[digit];
}
```

X86-64

```
# %rdi = z
# %rsi = digit
movl (%rdi,%rsi,4), %eax # z[digit]
```

- Register %rdi contains starting address of array
- Register %rsi contains array index
- Desired digit at %rdi + 4*%rsi
- Use memoryreference (%rdi,%rsi,4)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

6

73

Array Loop Example

```
void zincr(zip_dig z) {
    size_t i;
    for (i = 0; i < ZLEN; i++)
        z[i]++;
}
```

```
# %rdi = z
movl $0, %eax # i = 0
jmp .L3 # goto middle
.L4: # loop:
    addl $1, (%rdi,%rax,4) # z[i]++
    addq $1, %rax # i++
.L3: # middle
    cmpq $4, %rax # i:4
    jbe .L4 # if <=, goto loop
ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

7

74

Multidimensional (Nested) Arrays

Declaration

- `T A[R][C];`
- 2D array of data type T
- Rows, Ccolumns
- Type Telement requires Kbytes


Array Size

- $R * C * \text{Kbytes}$

Arrangement

- Row-Major Ordering

```
int A[R][C];
```



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

8


75

Nested Array Element Access

Array Elements

- $A[i][j]$ is element of type T, which requires Kbytes
- Address $A + i * (C * K) + j * K = A + (i * C + j) * K$

```
int A[R][C];
```




Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

12

76

Nested Array Element Access Code



```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

```
leaq    (%rdi,%rdi,4), %rax    # 5*index
addl    %rax, %rsi             # 5*index+dig
movl    pgh(,%rsi,4), %eax     # M[pgh + 4*(5*index+dig)]
```

Array Elements

- $pgh[index][dig]$ is `int`
- Address: $pgh + 20*index + 4*dig$
- $= pgh + 4*(5*index + dig)$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

13


77

Multi-Level Array Example

```
zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3
int *univ[UCOUNT] = { mit, cmu, ucb };
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
- 8 bytes
- Each pointer points to array of `int`'s




Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

14

78

Element Access in Multi-Level Array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi              # 4*digit
addq    univ(,%rdi,8), %rsi    # p = univ[index] + 4*digit
movl    (%rsi), %eax           # return *p
ret
```

Computation

- Element access $Mem[Mem[univ+8*index]+4*digit]$
- Must do two memory reads
 - First get pointer to row array
 - Then access element within array


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

15

79

Structure Representation

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



- Structure represented as block of memory
 - Big enough to hold all of the fields
- Fields ordered according to declaration
 - Even if another ordering could yield a more compact representation
- Compiler determines overall size + positions of fields
 - Machine-level program has no understanding of the structures in the source code


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

21

80

Structures & Alignment


Unaligned Data



```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

Aligned Data

- Primitive data type requires Kbytes
- Address must be multiple of K



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

24

81

Alignment Principles

Aligned Data

- Primitive data type requires Kbytes
- Address must be multiple of K
- Required on some machines; advised on x86-64

Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory trickier when datum spans 2 pages

Compiler

- Inserts gaps in structure to ensure correct alignment of fields

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

25

82

Specific Cases of Alignment (x86-64)

- 1 byte: char, ...
 - no restrictions on address
- 2 bytes: short, ...
 - lowest 1 bit of address must be 0₂
- 4 bytes: int, float, ...
 - lowest 2 bits of address must be 00₂
- 8 bytes: double, long, char *, ...
 - lowest 3 bits of address must be 000₂
- 16 bytes: long double (GCC on Linux)
 - lowest 4 bits of address must be 0000₂

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

83

Satisfying Alignment with Structures

Within structure:


- Must satisfy each element's alignment requirement

Overall structure placement

- Each structure has alignment requirement K
 - K = Largest alignment of any element
- Initial address & structure length must be multiples of K

Example:

- K=8, due to double element



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


27

84

Meeting Overall Alignment Requirement

- For largest alignment requirement K
- Overall structure must be multiple of K

```
struct S2 {
    double v;
    int i[2];
    char c;
} *p;
```



Multiple of K=8

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

28

85

Saving Space



- Put large data types first

```
struct S4 {
    char
    c;
    int i;
    char
    d;
```



```
struct S5 {
    int i;
    char
    c;
    char
    d;
} *p;
```

- Effect (K=4)



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

31

86

Machine-Level Programming V: Advanced Topics

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

1

87

Today

- Memory Layout
- Buffer Overflow
 - Vulnerability
 - Protection
- Unions

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

2

88

x86-64 Linux Memory Layout

not drawn to scale

Stack

- Runtime stack (8MB limit)
- E.g., local variables

Heap


- Dynamically allocated as needed
- When call malloc(), calloc(), new()

Data

- Statically allocated data
- E.g., global vars, static vars, string constants

Text / Shared Libraries

- Executable machine instructions
- Read-only

Hex Address → 400000
000000

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


89

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;
```

```
fun(0) □ 3.14
fun(1) □ 3.14
fun(2) □ 3.1399998664856
fun(3) □ 2.00000061035156
fun(4) □ 3.14
fun(6) □ Segmentation fault
```

Explanation:



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

90

Such problems are a BIG deal

Generally called a "bufferoverflow"

- when exceeding the memory size allocated for an array

Why a big deal?

- It's the #1 technical cause of security vulnerabilities
 - #1 overall cause is social engineering / user ignorance

Most common form

- Unchecked lengths on string inputs
- Particularly for bounded character arrays on the stack
 - sometimes referred to as stack smashing

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



91

Buffer Overflow Disassembly

echo:

```
00000000004006cf <echo>:
4006cf: 48 83 ec 18      sub    $0x18,%rsp
4006d3: 48 89 e7         mov    %rsp,%rdi
4006d6: e8 a5 ff ff ff   callq 400680 <get s>
4006db: 48 89 e7         mov    %rsp,%rdi
4006de: e8 3d fe ff ff   callq 400520 <put s@plt>
4006e3: 48 83 c4 18      add    $0x18,%rsp
4006e7: c3              retq
```

call_echo:

```
4006e8: 48 83 ec 08      sub    $0x8,%rsp
4006ec: b8 00 00 00 00   mov    $0x0,%eax
4006f1: e8 d9 ff ff ff   callq 4006cf <echo>
4006f6: 48 83 c4 08      add    $0x8,%rsp
4006fa: c3              retq
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



92

Buffer Overflow Stack Example#1

After call to gets

Stack Frame for call_echo			
00	00	00	00
00	40	y06	f6
00	32	31	30
39	38	37	36
35	34	33	32
31	30	29	28
27	26	25	24
23	22	21	20

```
void echo()
{
    char buf[4];
    gets(buf);
    ...
}

echo:
    subq    $24, %rsp
    movq    %rsp, %rdi
    call    gets
    ...
```

call_echo:

```
...
4006f1: callq 4006cf <echo>
4006f6: add    $0x8,%rsp
...
```

buf ← %rsp

```
unix> ./bufdemo-nsp
Type a string: 01234567890123456789012
01234567890123456789012
```

Overflowed buffer, but did not corrupt state

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



93

Buffer Overflow Stack Example#2

After call to gets

Stack Frame for call_echo			
00	00	00	00
00	40	y00	34
33	32	31	30
39	38	37	36
35	34	33	32
31	30	29	28
27	26	25	24
23	22	21	20

```
void echo()
{
    char buf[4];
    gets(buf);
    ...
}

echo:
    subq    $24, %rsp
    movq    %rsp, %rdi
    call    gets
    ...
```

call_echo:

```
...
4006f1: callq 4006cf <echo>
4006f6: add    $0x8,%rsp
...
```

buf ← %rsp

```
unix> ./bufdemo-nsp
Type a string: 0123456789012345678901234
Segmentation Fault
```

Overflowed buffer and corrupted return pointer

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



94

Buffer Overflow Stack Example#3

After call to gets

Stack Frame for call_echo			
00	00	00	00
00	40	y06	00
33	32	31	30
39	38	37	36
35	34	33	32
31	30	29	28
27	26	25	24
23	22	21	20

```
void echo()
{
    char buf[4];
    gets(buf);
    ...
}

echo:
    subq    $24, %rsp
    movq    %rsp, %rdi
    call    gets
    ...
```

call_echo:

```
...
4006f1: callq 4006cf <echo>
4006f6: add    $0x8,%rsp
...
```

buf ← %rsp

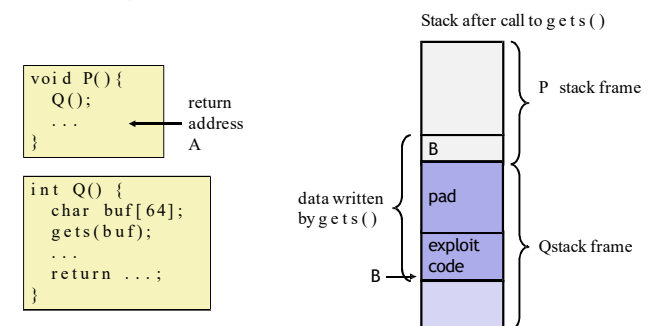
```
unix> ./bufdemo-nsp
Type a string: 012345678901234567890123
012345678901234567890123
```

Overflowed buffer, corrupted return pointer, but program seems to work!
"Returns" to unrelated code
Lots of things happen, without modifying critical state. Eventually executes `retq` back to main



95

Code Injection Attacks



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When Q executes `ret`, will jump to exploit code

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



96

1. Avoid Overflow Vulnerabilities in Code (!)

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    fgets(buf, 4, stdin);
    puts(buf);
}
```

For example, use library routines that limit string lengths

- `fgets` instead of `gets`
- `strncpy` instead of `strcpy`
- Don't use `scanf` with `%s` conversion specification
 - Use `fgets` to read the string
 - Or use `%ns` where `n` is a suitable integer

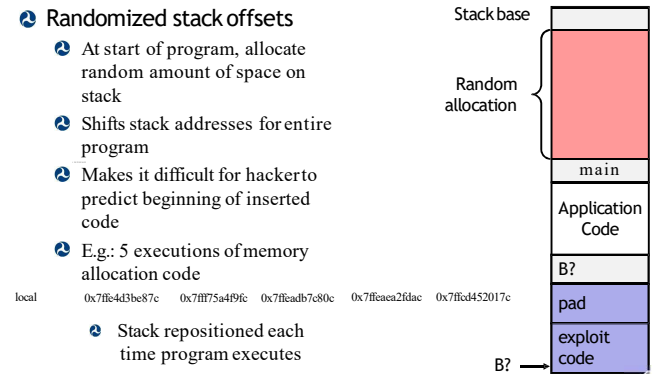
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

27

97

2. System-Level Protections can help

- Randomized stack offsets
 - At start of program, allocate random amount of space on stack
 - Shifts stack addresses for entire program
 - Makes it difficult for hacker to predict beginning of inserted code
 - E.g.: 5 executions of memory allocation code
- Stack repositioned each time program executes



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


28

98

2. System-Level Protections can help

Nonexecutable code segments

- In traditional x86, can mark region of memory as either "read-only" or "writable"
- Can execute anything readable
- X86-64 added explicit "execute" permission
- Stack marked as non-executable



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

29

99

3. Stack Canaries can help

Idea

- Place special value ("canary") on stack just beyond buffer
- Check for corruption before exiting function

GCC Implementation

- `-fstack-protector`
- Now the default (disabled earlier)

```
unix> ./bufdemo-sp
Type a string: 0123456
0123456
```

```
unix> ./bufdemo-sp
Type a string: 01234567
*** stack smashing detected ***
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

100

Protected Buffer Disassembly

echo:

```

40072f: sub    $0x18,%rsp
400733: mov    %fs:0x28,%rax
40073c: mov    %rax,0x8(%rsp)
400741: xor    %eax,%eax
400743: mov    %rsp,%rdi
400746: callq  4006e0 <gets>
40074b: mov    %rsp,%rdi
40074e: callq  400570 <puts@plt>
400753: mov    0x8(%rsp),%rax
400758: xor    %fs:0x28,%rax
400761: je     400768 <echo+0x39>
400763: callq  400580 <__stack_chk_fail@plt>
400768: add    $0x18,%rsp
40076c: retq

```


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

31

101

Setting Up Canary

Before call to gets



```

/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}

```

```

movq    %fs:40, %rax # Get canary
movq    %rax, 8(%rsp) # Place on stack
xorl    %eax, %eax # Erase canary

```

```

echo:
    movq    8(%rsp), %rax # Retrieve from stack
    xorq    %fs:40, %rax # Compare to canary
    je     .L6 # If same, OK
    call    __stack_chk_fail # FAIL

```

102

Return-Oriented Programming Attacks

- ⌚ Challenge (for hackers)
 - ⌚ Stack randomization makes it hard to predict buffer location
 - ⌚ Marking stack nonexecutable makes it hard to insert binary code
- ⌚ Alternative Strategy
 - ⌚ Use existing code
 - ⌚ E.g., library code from stdlib
 - ⌚ String together fragments to achieve overall desired outcome
 - ⌚ Does not overcome stack canaries
- ⌚ Construct program from gadgets
 - ⌚ Sequence of instructions ending in `ret`
 - ⌚ Encoded by single byte `0xc3`
 - ⌚ Code positions fixed from run to run
 - ⌚ Code is executable

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

34

103

Gadget Example #2

```

void setval(unsigned *p) {
    *p = 3347663060u;
}

```

```

<setval>:
4004d9: c7 07 d4 8 89 c7    movl $0xc78948d4, (%rdi)
4004df: c3                retq

```

Encodes `movq %rax, %rdi`

`rdi = rax`
Gadget address = `0x4004dc`


- ⌚ Repurpose byte codes

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

36

104

ROP Execution



- ⌚ Trigger with `ret` instruction
 - ⌚ Will start executing Gadget 1
- ⌚ Final `ret` in each gadget will start next one

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



37



105

Union Allocation

- ⌚ Allocate according to largest element
- ⌚ Can only use one field at a time

```
union U1 {
    char c;
    int i[2];
    double v;
} *up;
```

```
struct S1 {
    char c;
    int i[2];
    double v;
} *sp;
```



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



39

106

Floating Point

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



1

107

Today: Floating Point

- ⌚ Background: Fractional binary numbers
- ⌚ IEEE floating point standard: Definition
- ⌚ Example and properties
- ⌚ Rounding, addition, multiplication
- ⌚ Floating point in C
- ⌚ Summary

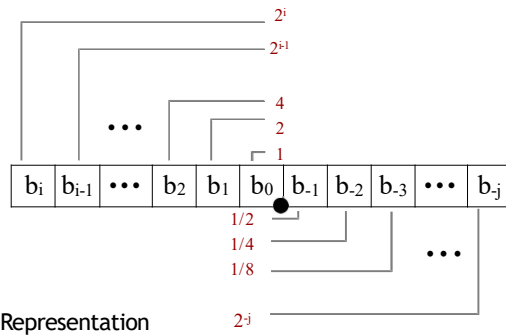
Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



2

108

Fractional Binary Numbers



Representation

- Bits to right of “binary point” represent fractional powers of 2
- Represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



109

Fractional Binary Numbers: Examples

- | Value | Representation |
|------------------|----------------|
| $5 \frac{3}{4}$ | 101.11_2 |
| $2 \frac{7}{8}$ | 10.111_2 |
| $1 \frac{7}{16}$ | 1.0111_2 |
- Observations
 - Divide by 2 by shifting right (unsigned)
 - Multiply by 2 by shifting left
 - Numbers of form $0.11111\dots_2$ are just below 1.0
 - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
 - Use notation $1.0 - \epsilon$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



110

Floating Point Representation

Numerical Form:

$$(-1)^s M 2^E$$

- Sign bit s determines whether number is negative or positive
- Significand M normally a fractional value in range $[1.0, 2.0)$.
- Exponent E weights value by power of two

Encoding

- MSBs is sign bit s
- exp field encodes E (but is not equal to E)
- frac field encodes M (but is not equal to M)




Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




111

Precision options


Single precision: 32 bits



Double precision: 64 bits



Extended precision: 80 bits (Intel only)



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



112

“Normalized” Values

$$v = (-1)^s M 2^E$$

- ⌚ When: $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$
- ⌚ Exponent coded as a biased value: $E = \text{Exp} - \text{Bias}$
 - ⌚ Exp: unsigned value of exp field
 - ⌚ Bias $= 2^{k-1} - 1$, where k is number of exponent bits
 - ⌚ Single precision: 127 (Exp: 1...254, E: -126...127)
 - ⌚ Double precision: 1023 (Exp: 1...2046, E: -1022...1023)
- ⌚ Significand coded with implied leading 1: $M = 1.\text{xxx}\dots\text{x}_2$
 - ⌚ xxx...x: bits of frac field
 - ⌚ Minimum when $\text{frac} = 000\dots 0$ ($M = 1.0$)
 - ⌚ Maximum when $\text{frac} = 111\dots 1$ ($M = 2.0 - \epsilon$)
 - ⌚ Get extra leading bit for “free”

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



113

Normalized Encoding Example

$$v = (-1)^s M 2^E$$

$$E = \text{Exp} - \text{Bias}$$

- ⌚ Value: float $F = 15213.0$;
 - ⌚ $15213_{10} = 11101101101101_2$
 $= 1.1101101101101_2 \times 2^{13}$
- ⌚ Significand
 - $M = 1.1101101101101_2$
 - $\text{frac} = 11011011011010000000000_2$
- ⌚ Exponent
 - $E = 13$
 - $\text{Bias} = 127$
 - $\text{Exp} = 140 = 10001100_2$
- ⌚ Result:

0

10001100

110110110110100000000000

s
exp
frac

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



114

Denormalized Values

$$v = (-1)^s M 2^E$$

$$E = 1 - \text{Bias}$$

- ⌚ Condition: $\text{exp} = 000\dots 0$
- ⌚ Exponent value: $E = 1 - \text{Bias}$ (instead of $E = 0 - \text{Bias}$)
- ⌚ Significand coded with implied leading 0: $M = 0.\text{xxx}\dots\text{x}_2$
 - ⌚ xxx...x: bits of frac
- ⌚ Cases
 - ⌚ $\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$
 - ⌚ Represents zero value
 - ⌚ Note distinct values: +0 and -0
 - ⌚ $\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$
 - ⌚ Numbers closest to 0.0
 - ⌚ Equispaced

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



115

Special Values

Carnegie Mellon

- ⌚ Condition: $\text{exp} = 111\dots 1$
- ⌚ Case: $\text{exp} = 111\dots 1, \text{frac} = 000\dots 0$
 - ⌚ Represents value \square (infinity)
 - ⌚ Operation that overflows
 - ⌚ Both positive and negative
 - ⌚ E.g., $1.0/0.0 = -1.0/-0.0 = +\square$, $1.0/-0.0 = -\square$
- ⌚ Case: $\text{exp} = 111\dots 1, \text{frac} \neq 000\dots 0$
 - ⌚ Not-a-Number (NaN)
 - ⌚ Represents case when no numeric value can be determined
 - ⌚ E.g., $\text{sqrt}(-1)$, $\square - \square$, \square / \square

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



116

Dynamic Range (Positive Only)

	s	exp	frac	E	Value	
Denormalized numbers	0	0000	000	-6	0	
	0	0000	001	-6	$1/8 * 1/64 = 1/512$	
	0	0000	010	-6	$2/8 * 1/64 = 2/512$	
	...					
	0	0000	110	-6	$6/8 * 1/64 = 6/512$	
Normalized numbers	0	0000	111	-6	$7/8 * 1/64 = 7/512$	largest denorm
	0	0001	000	-6	$8/8 * 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 * 1/64 = 9/512$	
	...					
	0	0110	110	-1	$14/8 * 1/2 = 14/16$	
	0	0110	111	-1	$15/8 * 1/2 = 15/16$	closest to 1 below
	0	0111	000	0	$8/8 * 1 = 1$	
	0	0111	001	0	$9/8 * 1 = 9/8$	closest to 1 above
	0	0111	010	0	$10/8 * 1 = 10/8$	
	...					
	0	1110	110	7	$14/8 * 128 = 224$	
	0	1110	111	7	$15/8 * 128 = 240$	largest norm
	0	1111	000	n/a	inf	

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

18

117

Special Properties of the IEEE Encoding

- FP Zero Same as Integer Zero
 - All bits = 0
- Can (Almost) Use Unsigned Integer Comparison
 - Must first compare sign bits
 - Must consider $-0 = 0$
 - NaNs problematic
 - Will be greater than any other values
 - What should comparison yield?
 - Otherwise OK
 - Denorm vs. normalized
 - Normalized vs. infinity

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

21

118

Floating Point Operations: Basic Idea

- $x +_f y = \text{Round}(x + y)$
- $x \square_f y = \text{Round}(x \square y)$
- Basic idea
 - First **compute exact result**
 - Make it fit into desired precision
 - Possibly overflow if exponent too large
 - Possibly **round to fit into frac**

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

23

119

Rounding Binary Numbers

- Binary Fractional Numbers
 - "Even" when least significant bit is 0
 - "Half way" when bits to right of rounding position = 100...
- Examples
 - Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded
2 3/32	10.00011 ₂	10.00 ₂	(<1/2—down)	2
2 3/16	10.00110 ₂	10.01 ₂	(>1/2—up)	2 1/4
2 7/8	10.11100 ₂	11.00 ₂	(1/2—up)	3
2 5/8	10.10100 ₂	10.10 ₂	(1/2—down)	2 1/2

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

120

FP Multiplication

$$(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$$

$$\text{Exact Result: } (-1)^s M 2^E$$

- Sign s: $s_1 \oplus s_2$
- Significand M: $M_1 \times M_2$
- Exponent E: $E_1 + E_2$

Fixing

- If $M \geq 2$, shift M right, increment E
- If E out of range, overflow
- Round M to fit for a c precision

Implementation

- Biggest chore is multiplying significands

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



27

121

Floating Point Addition


$$(-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$$

- Assume $E_1 > E_2$

Get binary points lined up

$$\text{Exact Result: } (-1)^s M 2^E$$

- Sign s, significand M:
 - Result of signed align & add
- Exponent E: E_1



Fixing

- If $M \geq 2$, shift M right, increment E
- If $M < 1$, shift M left k positions, decrement E by k
- Overflow if E out of range
- Round M to fit for a c precision

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



28

122

Floating Point in C

C Guarantees Two Levels

- `float` single precision
- `double` double precision

Conversions/Casting

- Casting between `int`, `float`, and `double` changes bit representation
- `double/float → int`
 - Truncates fractional part
 - Like rounding toward zero
 - Not defined when out of range or NaN: Generally sets to TMin
- `int → double`
 - Exact conversion, as long as `int` has ≤ 53 bit word size
- `int → float`
 - Will round according to rounding mode

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



30

123

The Memory Hierarchy

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



1

124

Today

- Storage technologies and trends
- Locality of reference
- Caching in the memory hierarchy

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



2

125

SRAM vs DRAM Summary

	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	4 or 6	1X	No	Maybe	100x	Cache memories
DRAM	1	10X	Yes	Yes	1X	Main memories, frame buffers

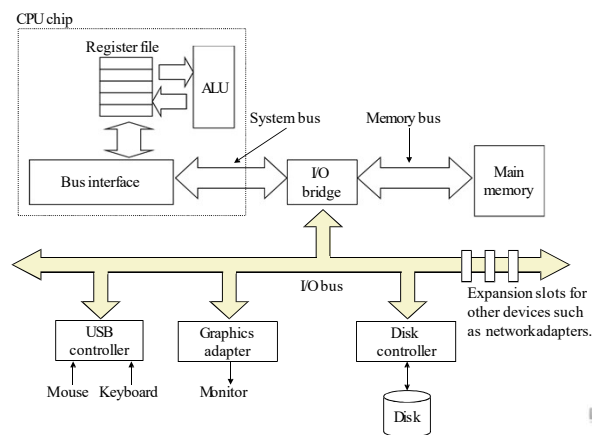
Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



4

126

I/O Bus



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



27

127

Locality Example

```
sum = 0;
for (i = 0; i < n; i++)
    sum += a[i];
return sum;
```

Data references

- Reference array elements in succession (stride-1 reference pattern).
- Reference variable `sum` each iteration.

Spatial locality

Temporal locality

Instruction references

- Reference instructions in sequence.
- Cycle through loop repeatedly.

Spatial locality

Temporal locality


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



38

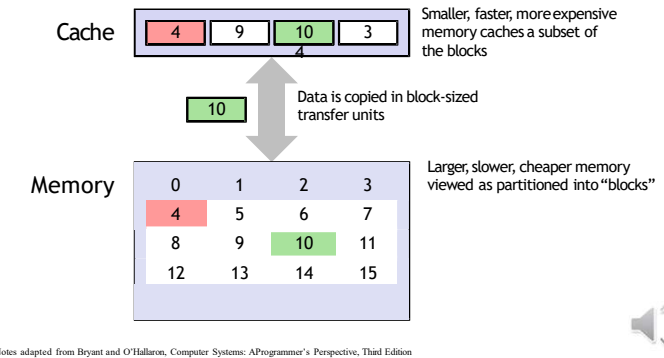
128

Example Memory Hierarchy



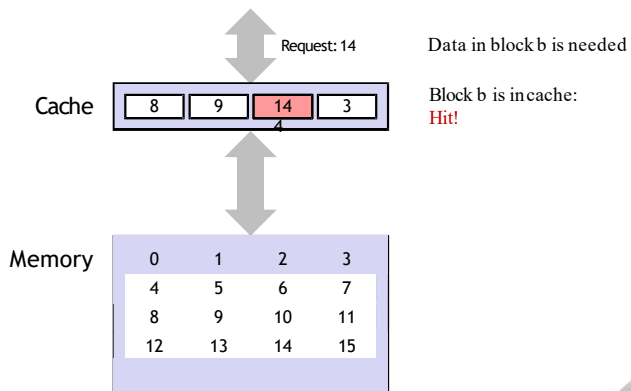
129

General Cache Concepts



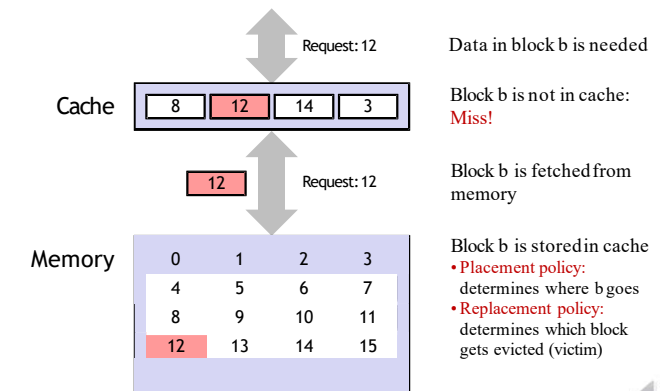
130

General Cache Concepts: Hit



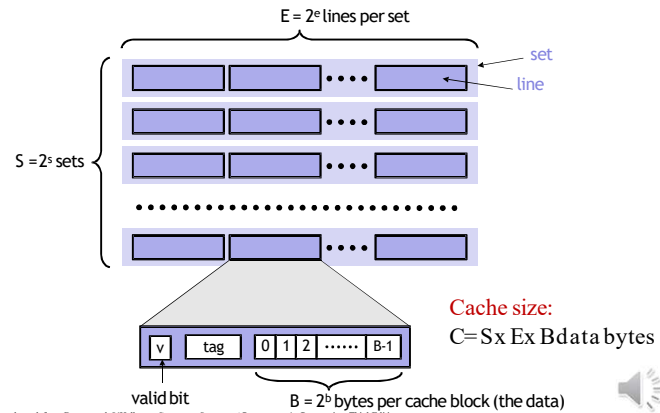
131

General Cache Concepts: Miss



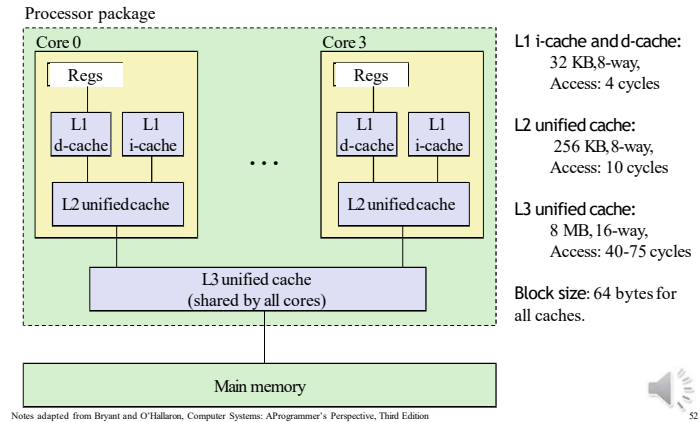
132

General Cache Organization (S, E, B)



133

Intel Core i7 Cache Hierarchy



134

Cache Performance Metrics

- Miss Rate
 - Fraction of memory references not found in cache (misses / accesses)
 - $= 1 - \text{hit rate}$
 - Typical numbers (in percentages):
 - 3-10% for L1
 - can be quite small (e.g., <1%) for L2, depending on size, etc.
- Hit Time
 - Time to deliver a line in the cache to the processor
 - includes time to determine whether the line is in the cache
 - Typical numbers:
 - 4 clock cycle for L1
 - 10 clock cycles for L2
- Miss Penalty
 - Additional time required because of a miss
 - typically 50-200 cycles for main memory (Trend: increasing!)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

135

Writing Cache Friendly Code

- Make the common case go fast
 - Focus on the inner loops of the core functions
- Minimize the misses in the inner loops
 - Repeated references to variables are good (temporal locality)
 - Stride-1 reference patterns are good (spatial locality)

Key idea: Our qualitative notion of locality is quantified through our understanding of cache memories

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

136

The Memory Mountain


- Read throughput (read bandwidth)
 - Number of bytes read from memory per second (MB/s)
- Memory mountain: Measured read throughput as a function of spatial and temporal locality.
 - Compact way to characterize memory system performance.

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



137

The Memory Mountain




58

138

Miss Rate Analysis for Matrix Multiply

- Assume:
 - Block size = 32B (big enough for four doubles)
 - Matrix dimension (N) is very large
 - Approximate $1/N$ as 0.0
 - Cache is not even big enough to hold multiple rows
- Analysis Method:
 - Look at access pattern of inner loop



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



139

Summary of Matrix Multiplication

```
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum=0.0;
    for (k=0; k<n; k++)
      sum+=a[i][k] * b[k][j];
    c[i][j] =sum;
  }
}
```

- ijk (& jik):
- 2 loads, 0 stores
 - misses/iter = 1.25

```
for (k=0; k<n; k++) {
  for (i=0; i<n; i++) {
    r =a[i][k];
    for (j=0; j<n; j++)
      c[i][j] +=r * b[k][j];
  }
}
```

- kij (& ikj):
- 2 loads, 1 store
 - misses/iter = 0.5

```
for (j=0; j<n; j++) {
  for (k=0; k<n; k++) {
    r =b[k][j];
    for (i=0; i<n; i++)
      c[i][j] +=a[i][k] * r;
  }
}
```

- jki (& kji):
- 2 loads, 1 store
 - misses/iter = 2.0

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




140

Blocked Matrix Multiplication

```
c=(double *) calloc(sizeof(double), n*n);

/* Multiply n x n matrices a and b */
void mmm(double*a, double*b, double*c, int n) {
    int i, j, k;
    for (i =0; i <n; i+=B) for
        for (j =0; j <n; j+=B)
            for (k =0; k <n; k+=B)
                /* B x B mini matrix multiplications */
                for (i1 =i; i1 <i+B; i1++)
                    for (j1 =j; j1 <j+B; j1++)
                        for (k1 =k; k1 <k+B; k1++)
                            c[i1*n+j1] +=a[i1*n +k1]*b[k1*n +j1];
}
matmult/bmm.c
```



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

141

Cache Miss Analysis


Assume:

- Cache block = 8 doubles
- Cache size $C \ll n$ (much smaller than n)
- Three blocks fit into cache: $3B^2 < C$

First (block) iteration:

- $B^2/8$ misses for each block
- $2n/B * B^2/8 = nB/4$ (omitting matrix c)

Afterwards in cache (schematic)



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

142

Blocking Summary

- No blocking: $(9/8) * n^3$
- Blocking: $1/(4B) * n^3$
- Suggest largest possible block size B , but limit $3B^2 < C$
- Reason for dramatic difference:
 - Matrix multiplication has inherent temporal locality:
 - Input data: $3n^2$, computation $2n^3$
 - Every array elements used $O(n)$ times!
 - But program has to be written properly

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

143

Program Optimization

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

144

Today

- Overview
- Generally Useful Optimizations
 - Codemotion/precomputation
 - Strength reduction
 - Sharing of common subexpressions
 - Removing unnecessary procedure calls
- Optimization Blockers
 - Procedure calls
 - Memory aliasing
- Exploiting Instruction Level Parallelism
- Dealing with Conditionals

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



145

Generally Useful Optimizations

- Optimizations that you or the compiler should do regardless of processor/ compiler
- Code Motion
 - Reduce frequency with which computation performed
 - If it will always produce same result
 - Especially moving code out of loop

```
void set_row(double *a, double *b,
             long i, long n)
{
    long j;
    for (j = 0; j < n; j++)
        a[n*i+j] = b[j];
}
```

```
long j;
int ni = n*i;
for (j = 0; j < n; j++)
    a[ni+j] = b[j];
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



146

Compiler-Generated Code Motion (-O1)

```
void set_row(double *a, double *b,
             long i, long n)
{
    long j;
    for (j = 0; j < n; j++)
        a[n*i+j] = b[j];
}
```

```
long j;
long ni = n*i;
double *rowp = a+ni;
for (j = 0; j < n; j++)
    *rowp++ = b[j];
```

```
set_row:
    testq    %rcx, %rcx          # Test n
    jle     .L1                 # If 0, goto done
    imulq    %rcx, %rdx          # ni = n*i
    leaq     (%rdi,%rdx,8), %rdx  # rowp = A+ ni*8
    movl     $0, %eax           # j = 0
    # loop:
    movsd    (%rsi,%rax,8), %xmm0 # t = b[j]
    movsd    %xmm0, (%rdx,%rax,8) # M[A+ni*8+j*8] = t
    addq     $1, %rax            # j++
    cmpq     %rcx, %rax          # j:n
    jne     .L3                 # if !=, goto loop
.L1:
    rep ; ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



147

Reduction in Strength

- Replace costly operation with simpler one
- Shift, add instead of multiply or divide
 - $16 * x \rightarrow x \ll 4$
 - Utility machine dependent
 - Depends on cost of multiply or divide instruction
 - On Intel Nehalem, integer multiply requires 3 CPU cycles
- Recognize sequence of products

```
for (i = 0; i < n; i++) {
    int ni = n*i;
    for (j = 0; j < n; j++)
        a[ni+j] = b[j];
    ni += n;
}
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



148

Share Common Subexpressions

- Reuse portions of expressions
- GC will do this with -O1

```
/* Sum neighbors of i,j */
up = val[(i-1)*n + j];
down = val[(i+1)*n + j];
left = val[i*n + j-1];
right = val[i*n + j+1];
sum = up + down + left + right;
```

3 multiplications: $i*n$, $(i-1)*n$, $(i+1)*n$

```
leaq 1(%rsi), %rax # i+1
leaq -1(%rsi), %r8 # i-1
imulq %rcx, %rsi # i*n
imulq %rcx, %rax # (i+1)*n
imulq %rcx, %r8 # (i-1)*n
addq %r8, %rsi # i*n+j
addq %rdx, %rax # (i+1)*n+j
addq %rdx, %r8 # (i-1)*n+j
```

```
long inj = i*n + j;
up = val[inj - n];
down = val[inj + n];
left = val[inj - 1];
right = val[inj + 1];
sum = up + down + left + right;
```

1 multiplication: $i*n$

```
imulq %rcx, %rsi # i*n
addq %rdx, %rsi # i*n+j
movq %rsi, %rax # i*n+j
subq %rcx, %rax # i*n+j-n
leaq (%rsi,%rcx), %rcx # i*n+j+n
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



149

Improving Performance

```
void lower(char *s)
{
    size_t i;
    size_t len = strlen(s);
    for (i = 0; i < len; i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```

- Move call to `strlen` outside of loop
- Since result does not change from one iteration to another
- Form of codemotion

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



150

Optimization Blocker: Procedure Calls

- Why couldn't compiler move `strlen` out of inner loop?

- Procedure may have side effects
 - Alters global state each time called
- Function may not return same value for given arguments
 - Depends on other parts of global state
- Procedure lower could interact with `strlen`

Warning:

- Compiler treats procedure call as a black box
- Weak optimizations near them

Remedies:

- Use of inline functions
 - GC does this with -O1
 - Within single file
- Do your own codemotion

```
size_t lencnt = 0;
size_t strlen(const char *s)
{
    size_t length = 0;
    while (*s != '\0') {
        s++; length++;
    }
    lencnt += length;
    return length;
}
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



151

Removing Aliasing

```
/* Sum rows is of n Xn matrix a
and store in vector b */
void sum_rows2(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        double val = 0;
        for (j = 0; j < n; j++)
            val += a[i*n + j];
        b[i] = val;
    }
}
```

```
# sum_rows2 inner loop
.L10:
    addsd (%rdi), %xmm0 # FP load + add
    addq $8, %rdi
    cmpq %rax, %rdi
    jnc .L10
```

- No need to store intermediate results

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



152

Exploiting Instruction-Level Parallelism

- Need general understanding of modern processor design
 - Hardware can execute multiple instructions in parallel
- Performance limited by data dependencies
- Simple transformations can yield dramatic performance improvement
 - Compilers often cannot make these transformations
 - Lack of associativity and distributivity in floating-point arithmetic

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



153

Basic Optimizations

```
void combine4(vec_ptr v, data_t *dest)
{
    long i;
    long length = vec_length(v);
    data_t *d = get_vec_start(v);
    data_t t = IDENT;
    for (i = 0; i < length; i++)
        t = t OP d[i];
    *dest = t;
}
```

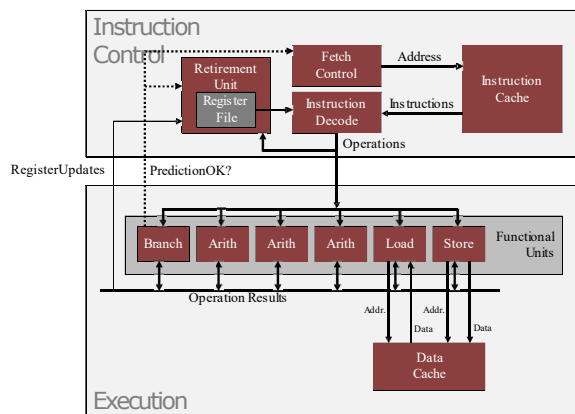
- Move vec_length out of loop
- Avoid bounds check on each cycle
- Accumulate in temporary

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



154

Modern CPU Design



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



155

Superscalar Processor

- **Definition:** A superscalar processor can issue and execute **multiple instructions in one cycle**. The instructions are retrieved from a sequential instruction stream and are usually scheduled dynamically.
- **Benefit:** without programming effort, superscalar processor can take advantage of the **instruction level parallelism** that most programs have
- Most modern CPUs are superscalar.
- Intel: since Pentium (1993)


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



156

Pipelined Functional Units

```
long mult_eg(long a, long b, long c) {
    long p1 = a*b;
    long p2 = a*c;
    long p3 = p1 * p2;
    return p3;
}
```



	Time						
	1	2	3	4	5	6	7
Stage 1	a*b	a*c			p1*p2		
Stage 2		a*b	a*c			p1*p2	
Stage 3			a*b	a*c			p1*p2

- Divide computation into stages
- Pass partial computations from stage to stage
- Stage i can start on new computation once values passed to i+1
- E.g., complete 3 multiplications in 7 cycles, even though each requires 3 cycles

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



157

Haswell CPU

- 8 Total Functional Units
 - Multiple instructions can execute in parallel
 - 2 load, with address computation
 - 1 store, with address computation
 - 4 integer
 - 2 FP multiply
 - 1 FP add
 - 1 FP divide
 - Some instructions take > 1 cycle, but can be pipelined
- | Instruction | Latency | Cycles/Issue |
|---------------------------|---------|--------------|
| Load / Store | 4 | 1 |
| Integer Multiply | 3 | 1 |
| Integer/Long Divide | 3-30 | 1 |
| Single/Double FP Multiply | 5 | 3 |
| Single/Double FP Add | 3 | 1 |
| Single/Double FP Divide | 3-15 | 1 |

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



158

Loop Unrolling with Reassociation (2x1a)

```
void unroll2aa_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = x OP d[i] OP d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OP d[i];
    }
    *dest = x;
}
```

Compare to before

```
x = (x OP d[i]) OP d[i+1];
```

- Can this change the result of the computation?
- Yes, for FP. Why?

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



159

Loop Unrolling with Separate Accumulators (2x2)

```
void unroll2a_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x0 = IDENT;
    data_t x1 = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x0 = x0 OP d[i];
        x1 = x1 OP d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x0 = x0 OP d[i];
    }
    *dest = x0 OP x1;
}
```

- Different form of reassociation

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



160

Unrolling & Accumulating

- Idea
 - Can unroll to any degree
 - Can accumulate results in parallel
 - Must be multiple of K
- Limitations
 - Diminishing returns
 - Cannot go beyond throughput limitations of execution units
 - Large overhead for short lengths
 - Finish off iterations sequentially


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



161

Branch Prediction

- Idea
 - Guess which way branch will go
 - Begin executing instructions at predicted position
 - But don't actually modify register or memory data

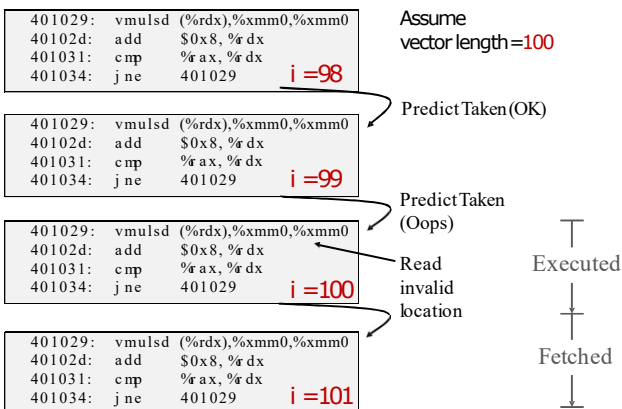


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



162

Branch Prediction Through Loop



- Cost of wrong prediction can be Multiple clock cycles on modern processor⁵³

Getting High Performance

- Good compiler and flags
- Don't do anything stupid
 - Watch out for hidden algorithmic inefficiencies
 - Write compiler-friendly code
 - Watch out for optimization blockers:
 - procedure calls & memory references
 - Look carefully at innermost loops (where most work is done)
- Tune code for machine
 - Exploit instruction-level parallelism
 - Avoid unpredictable branches
 - Make code cache friendly (Covered later in course)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



164

Domain Decomposition

First, decide how data elements should be divided among processors

Second, decide which tasks each processor should be doing


Example: Vector addition



165

Domain Decomposition

Find the largest element of an array



166

Task (Functional) Decomposition

First, divide tasks among processors


Second, decide which data elements are going to be accessed (read and/or written) by which processors

Example: Event-handler for GUI

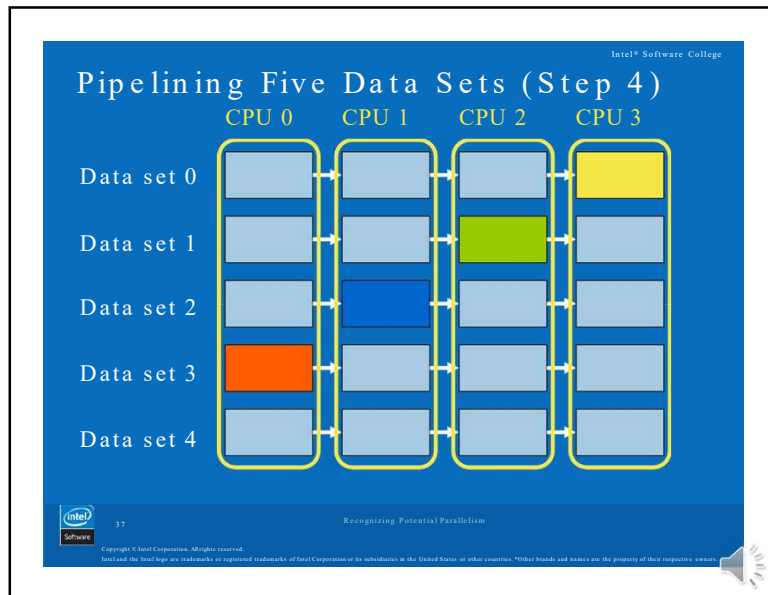


167

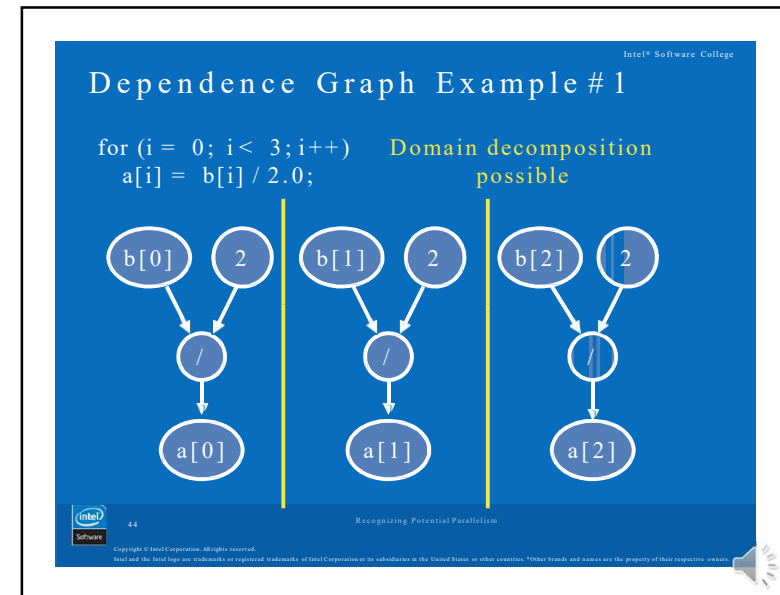
Task Decomposition



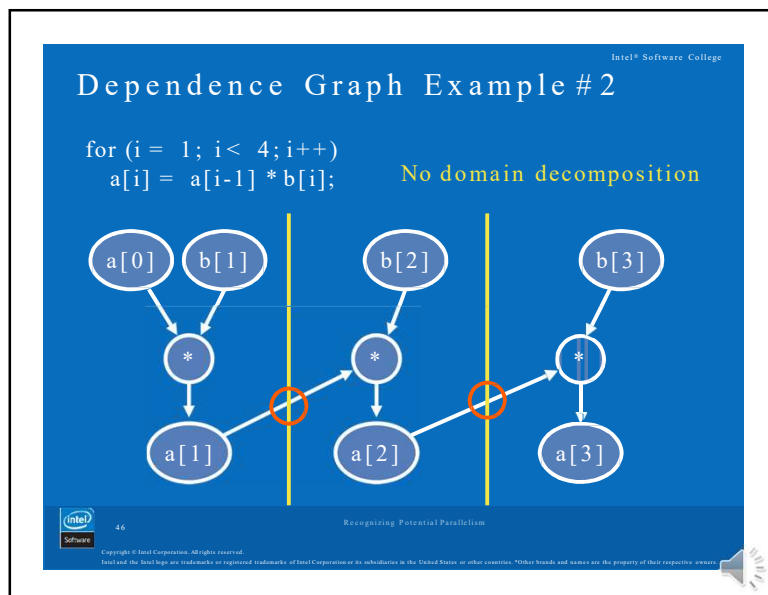
168




169



170




171



172

Relating Fork/Join to Code




Intel Software College

55 Shared-Memory Model and Threads

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

173

Domain Decomposition Using Threads




Intel Software College

58 Shared-Memory Model and Threads

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

174

Functional Decomposition Using Threads




Intel Software College

59 Shared-Memory Model and Threads

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

175

Pipelining Using Threads




Intel Software College

60 Shared-Memory Model and Threads

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

176

The Threads Model



177

Pragma: parallel for

The compiler directive

```
#pragma omp parallel for
```

tells the compiler that the for loop which immediately follows can be executed in parallel

The number of loop iterations must be computable at run time before loop executes

Loop must not contain a break, return, or exit

Loop must not contain a goto to a label outside loop

178

Matching Threads with CPUs (cont.)

Function `omp_set_num_threads` allows you to set the number of threads that should be active in parallel sections of code

```
void omp_set_num_threads (int t);
```

The function can be called with different arguments at different points in the program

Example:

```
int t;
```

```
...
```

```
omp_set_num_threads (t);
```

179

Problem Solved with private Clause

```
main () {
  int i, j, k;
  float **a, **b;

  ...
  for (k = 0; k < N; k++)
    #pragma omp parallel for private (j)
    for (i = 0; i < N; i++)
      for (j = 0; j < N; j++)
        a[i][j] = MIN(a[i][j], a[i][k] + a[k][j]);
}
```

Tells compiler to make listed variables private

180

Solution

```
int i;
float *a, *b, *c, tmp;
...
#pragma omp parallel for private (tmp)
for (i = 0; i < N; i++) {
    tmp = a[i] / b[i];
    c[i] = tmp * tmp;
}
```



181

Clause: firstprivate

The `firstprivate` clause tells the compiler that the private variable should inherit the value of the shared variable upon loop entry

The value is assigned once per thread, not once per loop iteration



182

Example

```
a[0] = 0.0;
for (i = 1; i < N; i++)
    a[i] = alpha (i, a[i-1]);
#pragma omp parallel for firstprivate (a)
for (i = 0; i < N; i++) {
    b[i] = beta (i, a[i]);
    a[i] = gamma (i);
    c[i] = delta (a[i], b[i]);
}
```



183

Clause: lastprivate

The `lastprivate` clause tells the compiler that the value of the private variable after the sequentially last loop iteration should be assigned to the shared variable upon loop exit

In other words, when the thread responsible for the sequentially last loop iteration exits the loop, its copy of the private variable is copied back to the shared variable



184

Example

```
#pragma omp parallel for lastprivate (x)
for (i = 0; i < N; i++) {
    x = foo (i);

    y[i] = bar(i, x);
}
last_x = x;
```



185

Pragma: parallel

In the effort to increase grain size, sometimes the code that should be executed in parallel goes beyond a single for loop

The `parallel` pragma is used when a block of code should be executed in parallel



186

Pragma: for

The `for` pragma is used inside a block of code already marked with the `parallel` pragma

It indicates a `for` loop whose iterations should be divided among the active threads

There is a barrier synchronization of the threads at the end of the `for` loop



187

Pragma: single

The `single` pragma is used inside a parallel block of code

It tells the compiler that only a single thread should execute the statement or block of code immediately following



188

Clause: nowait

The `nowait` clause tells the compiler that there is no need for a barrier synchronization at the end of a parallel `for` loop or single block of code



189


Solution: parallel, for, single Pragma

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i = 0; i < N; i++)
        a[i] = alpha(i);
    #pragma omp single nowait
    if (delta < 0.0) printf ("delta < 0.0\n");
    #pragma omp for
    for (i = 0; i < N; i++)
        b[i] = beta(i, delta);
}
```



190

Another Timing \Rightarrow Incorrect Sum



191

Mutual Exclusion

We can prevent the race conditions described earlier by ensuring that only one thread at a time references and updates shared variable or data structure

Mutual exclusion refers to a kind of synchronization that allows only a single thread or process at a time to have access to a shared resource

Mutual exclusion is implemented using some form of locking



192

Locking Mechanism

The previous method failed because checking the value of flag and setting its value were two distinct operations

We need some sort of atomic test-and-set

Operating system provides functions to do this

The generic term “lock” refers to a synchronization mechanism used to control access to shared resources



109

Conquering Race Conditions

Copyright © Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.



193

Critical Sections

A critical section is a portion of code that threads execute in a mutually exclusive fashion

The `critical` pragma in OpenMP immediately precedes a statement or block representing a critical section

Good news: critical sections eliminate race conditions

Bad news: critical sections are executed sequentially

More bad news: you have to identify critical sections yourself



110

Conquering Race Conditions

Copyright © Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.



194

OpenMP reduction Clause

Reductions are so common that OpenMP provides a reduction clause for the `parallel for` pragma

Eliminates need for

Creating private variable

Dividing computation into accumulation of local answers that contribute to global result



116

Conquering Race Conditions

Copyright © Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.



195

Solution #4

```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for private(x) \
                        reduction(+: area)
for (i = 0; i < n; i++) {
    x = (i + 0.5)/n;
    area += 4.0/(1.0 + x*x);
}
pi = area / n;
```



117

Conquering Race Conditions

Copyright © Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.



196

Important: Lock Data, Not Code

- Locks should be associated with data objects
- Different data objects should have different locks
- Suppose lock associated with critical section of code instead of data object
- Mutual exclusion can be lost if same object manipulated by two different functions
- Performance can be lost if two threads manipulating different objects attempt to execute same function




197

Deadlock

A situation involving two or more threads (processes) in which no thread may proceed because each is waiting for a resource held by another

Can be represented by a resource allocation graph



A graph of deadlock contains a cycle



198

More on Deadlocks


- A program exhibits a global deadlock if every thread is blocked
- A program exhibits local deadlock if only some of the threads in the program are blocked
- A deadlock is another example of a nondeterministic behavior exhibited by a parallel program
- Adding debugging output to detect source of deadlock can change timing and reduce chance of deadlock occurring



199

Deadlock Prevention Strategies

Don't allow mutually exclusive access to resource	Make resource shareable
Don't allow threads to wait while holding resources	Only request resources when have none. That means only hold one resource at a time or request all resources at once.
Allow resources to be taken away from threads.	Allow preemption. Works for CPU and memory. Doesn't work for locks.
Ensure no cycle in request allocation graph.	Rank resources. Threads must acquire resources in order.



200

Work Pool Analogy



More rows than workers

Each worker takes an unpicked row and picks the crop

After completing a row, the worker takes another unpicked row

Process continues until all rows have been harvested



140

Implementing Task Decompositions

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

201

Problem Site

```
int main ()
{
    struct board *stack;
    ...
    #pragma omp parallel
    search_for_solutions
        (n, stack, &num_solutions);
    ...
}

void search_for_solutions (int n,
    struct board *stack, int *num_solutions)
{
    ...
    while (stack != NULL) ...
}
```




153

Implementing Task Decompositions

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

202

Remedy 2: Use Indirection



159

Implementing Task Decompositions

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

203

Corrected Stack Access Function

```
void search_for_solutions (int n,
    struct board **stack, int *num_solutions)
{
    struct board *ptr;
    void search (int, struct board *, int *);

    while (*stack != NULL) {
        #pragma omp critical
        { ptr = *stack;
          *stack = (*stack)->next; }
        search (n, ptr, num_solutions);
        free (ptr);
    }
}
```



161

Implementing Task Decompositions

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

204

Parallel Sections

```
#pragma omp parallel sections
{
  <code block A>
  #pragma omp section
  <code block B>
  #pragma omp section
  <code block C>
}
```

Meaning: The following block contains sub-blocks that may execute in parallel

Each block executed by one thread

Dividers between sections



169

Implementing Task Decompositions

Copyright © Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. Other brands and names are the property of their respective owners.

205

RISC vs. CISC Machines

Feature	RISC	CISC
Registers	32	6, 8, 16
Register Classes	One	Some
Arithmetic Operands	Registers	Memory+Registers
Instructions	3-addr	2-addr
Addressing Modes	r M[r+c] (l,s)	several
Instruction Length	32 bits	Variable

Notes adapted from various sources



206

MIPS is a Load-Store Architecture

- Every operand of a MIPS instruction must be in a register (with some exceptions)
- Variables must be loaded into registers
- Results have to be stored back into memory
- Example C fragment...


```
a = b + c;
d = a + b;
```
- ... would be "translated" into something like:


```
Load b into register Rx
Load c into register Ry
Rz <- Rx + Ry
Store Rz into a
Rz <- Rz + Rx
Store Rz into d
```

Notes adapted from various sources



207

MIPS Register Names and Conventions

Register	Name	Function	Comment
\$0	zero	Always 0	No-op on write
\$1	\$at	reserved for assembler	don't use it!
\$2-3	\$v0-v1	expression eval./function return	
\$4-7	\$a0-a3	proc/funct call parameters	
\$8-15	\$t0-t7	volatile temporaries	not saved on call
\$16-23	\$s0-s7	temporaries (saved across calls)	saved on call
\$24-25	\$t8-t9	volatile temporaries	not saved on call
\$26-27	\$k0-k1	reserved kernel/OS	don't use them
\$28	\$gp	pointer to global data area	
\$29	\$sp	stack pointer	
\$30	\$fp	frame pointer	
\$31	\$ra	proc/funct return address	

Notes adapted from various sources



208

MIPS Instruction Types

- As we said earlier, there are very few basic operations :
 1. Memory access (load and store)
 2. Arithmetic (addition, subtraction, etc)
 3. Logical (and, or, xor, etc)
 4. Comparison (less-than, greater-than, etc)
 5. Control (branches, jumps, etc)
- We'll use the following notation when describing instructions:

rd: destination register (modified by instruction)

rs: source register (read by instruction)

rt: source/destination register (read or read+modified)

immed: a 16-bit value

Notes adapted from various sources



209

Load and Store Examples

- Load a word from memory:

```
lw rt, offset(base) # rt <- memory[base+offset]
```

- Store a word into memory:

```
sw rt, offset(base) # memory[base+offset] <- rt
```

- For smaller units (bytes, half-words) only the lower bits of a register are accessible. Also, for loads, you need to specify whether to sign or zero extend the data.

```
lb rt, offset(base) # rt <- sign-extended byte
```

```
lbu rt, offset(base) # rt <- zero-extended byte
```

```
sb rt, offset(base) # store low order byte of rt
```

Notes adapted from various sources



210

Arithmetic Instructions

Opcode	Operands	Comments
ADD	rd, rs, rt	# rd <- rs + rt
ADDI	rt, rs, immed	# rt <- rs + immed
SUB	rd, rs, rt	# rd <- rs - rt

Examples:

```
ADD    $8, $8, $10    # r8 <- r9 + r10
ADD    $t0, $t1, $t2  # t0 <- t1 + t2
SUB    $s0, $s0, $s1  # s0 <- s0 - s1
ADDI   $t3, $t4, 5    # t3 <- t4 + 5
```

Notes adapted from various sources



211

Flow of Control: Conditional Branches

```
BEQ    rs, rt, target # branch if rs == rt
BNE    rs, rt, target # branch if rs != rt
```

Comparison Between Registers

- What if you want to branch if R6 is greater than R7?
- We can use the SLT instruction:

```
SLT    rd, rs, rt    # if rs < rt then rd <- 1
                        # else rd <- 0
SLTU   rd, rs, rt    # same, but rs,rt unsigned
```

- Example: Branch to L1 if \$5 > \$6

```
SLT    $7, $6, $5    # $7 = 1, if $6 < $5
BNE    $7, $0, L1
```



212

Jump Instructions

- Jump instructions allow for unconditional transfer of control:

```
J      target    # go to specified target
JR     rs        # jump to addr stored in rs
```

- Jump and link is used for procedure calls:

```
JAL    target    # jump to target, $31 <- PC
JALR   rs, rd     # jump to addr in rs
                        # rd <- PC
```

- When calling a procedure, use JAL; to return, use JR \$31

Notes adapted from various sources

11

213

Logic Instructions

- Used to manipulate bits within words, set up masks, etc.

Opcode	Operands	Comments
AND	rd, rs, rt	# rd <- AND(rs, rt)
ANDI	rt, rs, immed	# rt <- AND(rs, immed)
OR	rd, rs, rt	
ORI	rt, rs, immed	
XOR	rd, rs, rt	
XORI	rt, rs, immed	

- The immediate constant is limited to 16 bits
- To load a constant in the 16 upper bits of a register we use LUI:

Opcode	Operands	Comments
LUI	rt, immed	# rt<31,16> <- immed
		# rt<15,0> <- 0

Notes adapted from various sources

12

214

Pseudoinstructions

Data moves

Name	Assembly syntax	Expansion	Operation in C
move	move \$t, \$s	addiu \$t, \$s, 0	t = s
clear	clear \$t	addu \$t, \$zero, \$zero	t = 0
load 16-bit immediate	li \$t, C	addiu \$t, \$zero, C_lo	t = C
load 32-bit immediate	li \$t, C	lui \$t, C_hi ori \$t, \$t, C_lo	t = C
load label address	la \$t, A	lui \$t, A_hi ori \$t, \$t, A_lo	t = A

Notes adapted from various sources

13

215

System Calls

Service	Code	Arguments	Result
print integer	1	\$a0=integer	Console print
print string	4	\$a0=string address	Console print
read integer	5		\$a0=result
read string	8	\$a0=string address \$a1=length limit	Console read
exit	10		end of program

Notes adapted from various sources

14

216

Hello World

```
.text      # text segment
.global __start

__start:   # execution starts here
    la $a0, str    # put string address into a0
    li $v0, 4      #
    syscall        # print
    li $v0, 10     #
    syscall        # au revoir...
.data      # data segment
str: .asciiz "hello world\n"
```



217

Virtual Memory

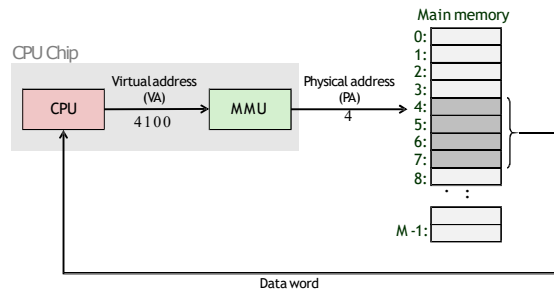
- Address spaces
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

1

218

A System Using Virtual Addressing



- Used in all modern servers, laptops, and smart phones
- One of the great ideas in computer science

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

3

219

Address Spaces

- **Linear address space:** Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots\}$
- **Virtual address space:** Set of $N=2^n$ virtual addresses
 $\{0, 1, 2, 3, \dots, N-1\}$
- **Physical address space:** Set of $M=2^m$ physical addresses
 $\{0, 1, 2, 3, \dots, M-1\}$

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

4

220

Why Virtual Memory (VM)?

- Uses main memory efficiently
 - Use DRAM as a cache for parts of a virtual address space
- Simplifies memory management
 - Each process gets the same uniform linear address space
- Isolates address spaces
 - One process can't interfere with another's memory
 - User program cannot access privileged kernel information and code

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

5

221

DRAM Cache Organization

- DRAM cache organization driven by the enormous miss penalty
 - DRAM is about **10x** slower than SRAM
 - Disk is about **10,000x** slower than DRAM
- Consequences
 - Large page (block) size: typically 4 KB, sometimes 4 MB
 - Fully associative
 - Any VP can be placed in any PP
 - Requires a "large" mapping function – different from cache memories
 - Highly sophisticated, expensive replacement algorithms
 - Too complicated and open-ended to be implemented in hardware
 - Write-back rather than write-through

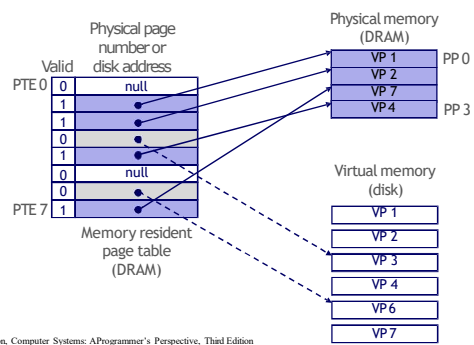
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

8

222

Enabling Data Structure: Page Table

- A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages.
- Per-process kernel data structure in DRAM



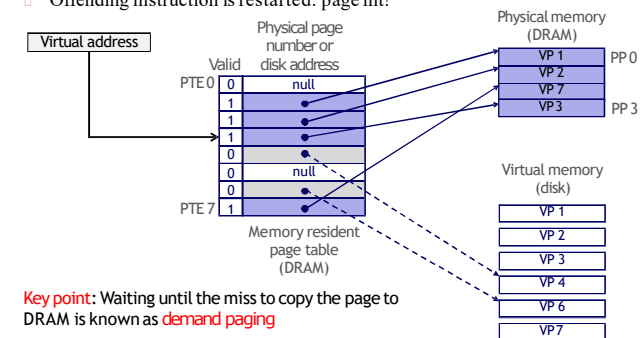
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

9

223

Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

15

224

Locality to the Rescue Again!

- Virtual memory seems terribly inefficient, but it works because of locality.
- At any point in time, programs tend to access a set of active virtual pages called the **working set**
 - Programs with better temporal locality will have smaller working sets
- If (working set size < main memory size)
 - Good performance for one process after compulsory misses
- If (SUM(working set sizes) > main memory size)
 - **Thrashing**: Performance meltdown where pages are swapped (copied) in and out continuously


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

17

225

VM as a Tool for Memory Management

- Key idea: each process has its own virtual address space
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory
 - Well-chosen mappings can improve locality




Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

19

226

VM as a Tool for Memory Management

- Simplifying memory allocation
 - Each virtual page can be mapped to any physical page
 - A virtual page can be stored in different physical pages at different times
- Sharing code and data among processes
 - Map virtual pages to the same physical page (here: PP 6)



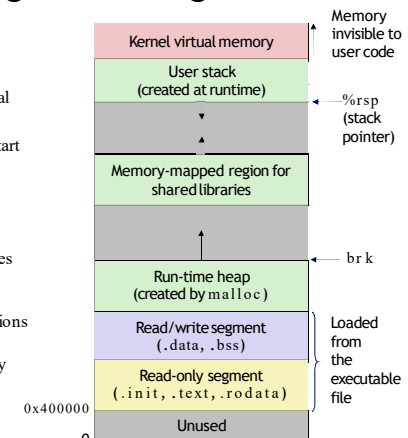
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

20

227

Simplifying Linking and Loading

- Linking
 - Each program has similar virtual address space
 - Code, data, and heap always start at the same addresses.
- Loading
 - `execve` allocates virtual pages for .text and .data sections & creates PTEs marked as invalid
 - The .text and .data sections are copied, page by page, on demand by the virtual memory system



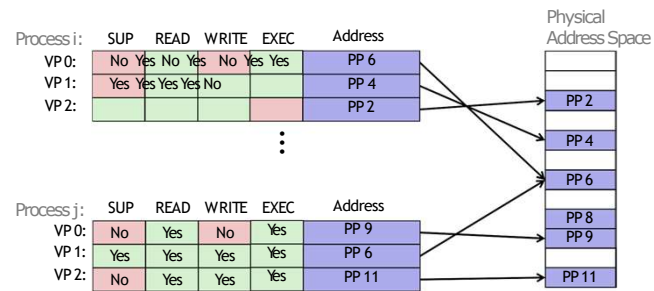
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

21

228

VM as a Tool for Memory Protection

- Extend PTEs with permission bits
- MMU checks these bits on each access



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

23

229

Summary of Address Translation Symbols

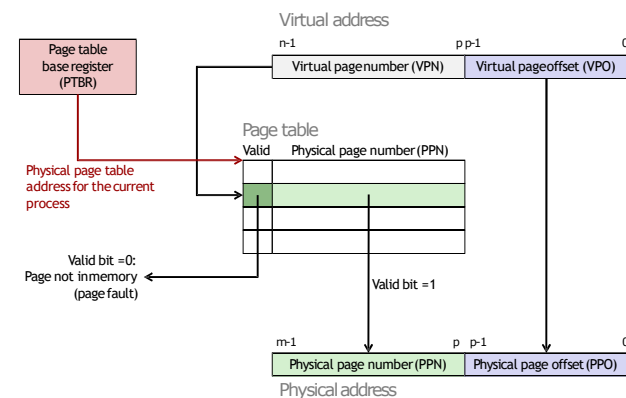
- Basic Parameters
 - $N = 2^n$: Number of addresses in virtual address space
 - $M = 2^m$: Number of addresses in physical address space
 - $P = 2^p$: Page size (bytes)
- Components of the virtual address (VA)
 - TLBI: TLB index
 - TLBT: TLB tag
 - VPO: Virtual page offset
 - VPN: Virtual page number
- Components of the physical address (PA)
 - PPO: Physical page offset (same as VPO)
 - PPN: Physical page number

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

230

Address Translation With a Page Table




Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

27

231

Address Translation: Page Hit




- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

28

232

Address Translation: Page Fault



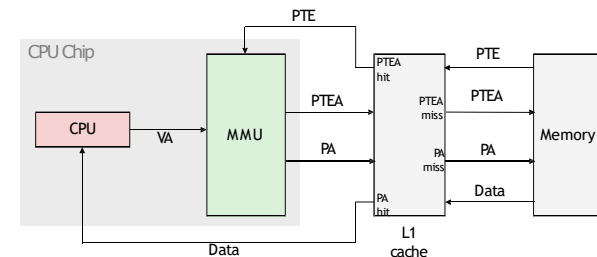
- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

29

233

Integrating VM and Cache



VA: virtual address, PA: physical address, PTE: page table entry, PTEA = PTE address

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

234

Speeding up Translation with a TLB

- Page table entries (PTEs) are cached in L1 like any other memory word
 - PTEs may be evicted by other data references
 - PTE hit still requires a small L1 delay
- Solution: **Translation Lookaside Buffer (TLB)**
 - Small set-associative hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages

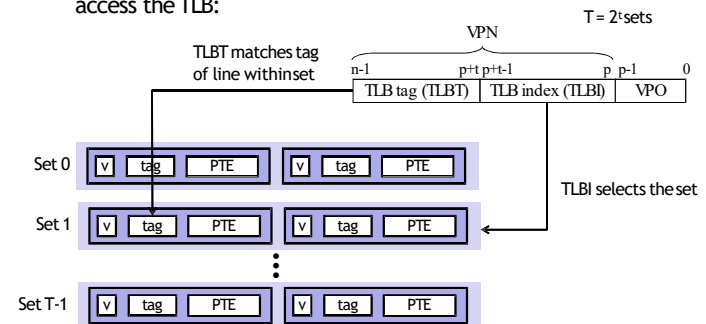
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

31

235

Accessing the TLB

- MMU uses the VPN portion of the virtual address to access the TLB:

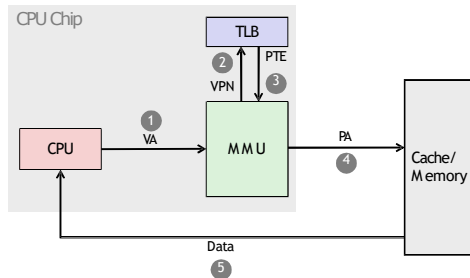


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

32

236

TLB Hit



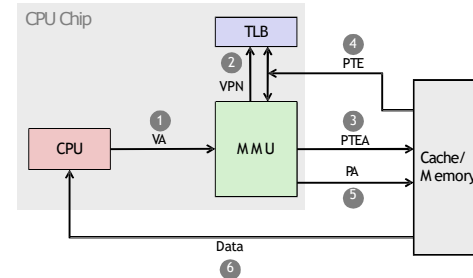
A TLB hit eliminates a memory access

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

33

237

TLB Miss



A TLB miss incurs an additional memory access (the PTE)
Fortunately, TLB misses are rare. Why?


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

34

238

Multi-Level Page Tables

- Suppose:
 - 4KB (2^{12}) page size, 48-bit address space, 8-byte PTE
- Problem:
 - Would need a 512 GB page table!
 - $2^{48} * 2^{-12} * 2^3 = 2^{39}$ bytes
- Common solution: Multi-level page table
- Example: 2-level page table
 - Level 1 table: each PTE points to a page table (always memory resident)
 - Level 2 table: each PTE points to a page (paged in and out like any other data)

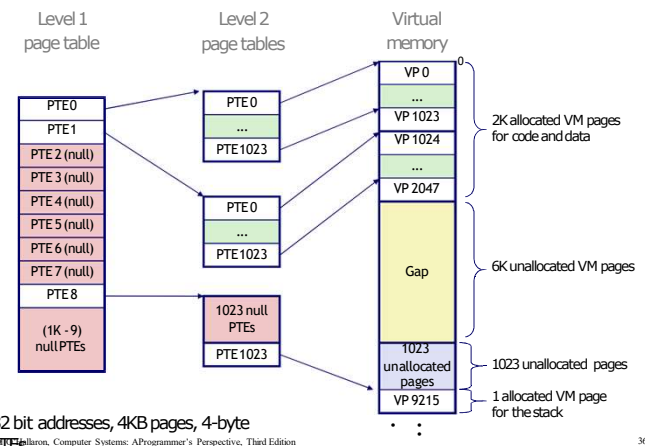


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

35

239

A Two-Level Page Table Hierarchy



32 bit addresses, 4KB pages, 4-byte PTEs

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

36

240

Summary

- Programmer's view of virtual memory
 - Each process has its own private linear address space
 - Cannot be corrupted by other processes
- System view of virtual memory
 - Uses memory efficiently by caching virtual memory pages
 - Efficient only because of locality
 - Simplifies memory management and programming
 - Simplifies protection by providing a convenient interpositioning point to check permissions

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

38

241

Altering the Control Flow

- Up to now: two mechanisms for changing control flow:
 - Jumps and branches
 - Call and return
 React to changes in **program state**
- Insufficient for a useful system:

Difficult to react to changes in **system state**

 - Data arrives from a disk or a network adapter
 - Instruction divides by zero
 - User hits Ctrl-C at the keyboard
 - System timer expires
- System needs mechanisms for “exceptional control flow”

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



242

Exceptional Control Flow

- Exists at all levels of a computer system
- Low level mechanisms
 - 1. **Exceptions**
 - Change in control flow in response to a system event (i.e., change in system state)
 - Implemented using combination of hardware and OS software
- Higher level mechanisms
 - 2. **Process context switch**
 - Implemented by OS software and hardware timer
 - 3. **Signals**
 - Implemented by OS software
 - 4. **Nonlocal jumps**: `setjmp()` and `longjmp()`
 - Implemented by Cruntime library


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



243

Exceptions

- An **exception** is a transfer of control to the OS kernel in response to some event (i.e., change in processor state)
 - Kernel is the memory-resident part of the OS
 - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C

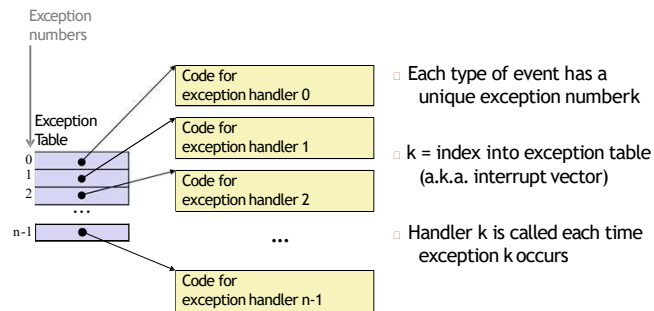


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



244

Exception Tables



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



7

245

Asynchronous Exceptions (Interrupts)

- Caused by events external to the processor
 - Indicated by setting the processor's interrupt pin
 - Handler returns to "next" instruction
- Examples:
 - Timer interrupt
 - Every few ms, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
 - I/O interrupt from external device
 - Hitting Ctrl-C at the keyboard
 - Arrival of a packet from a network
 - Arrival of data from a disk

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



8

246

Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:
 - Traps**
 - Intentional
 - Examples: system calls, breakpoint traps, special instructions
 - Returns control to "next" instruction
 - Faults**
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
 - Either re-executes faulting ("current") instruction or aborts
 - Aborts**
 - Unintentional and unrecoverable
 - Examples: illegal instruction, parity error, machine check
 - Aborts current program

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



9

247

System Calls

- Each x86-64 system call has a unique ID number
- Examples:

Number	Name	Description
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




10

248

System Call Example: OpeningFile

- User calls: `open(filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

```
0000000000e5d70<__open>:
...
e5d79: b8 02 00 00    mov $0x2,%eax # open is syscall #2
e5d7e: 0f 05          syscall      # Return value in %rax
e5d80: 48 3d 01 f0 ff  cmp $0xfffffffff001,%rax
...
e5dfa: c3            retq
```



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

11

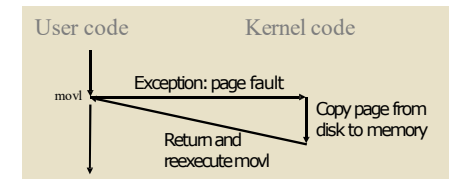
249

Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];
main ()
{
    a[500] = 13;
}
```

```
80483b7: c7 05 10 9d 04 0d movl $0xd,0x8049d10
```



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


12

250

Fault Example: Invalid Memory Reference

```
int a[1000];
main ()
{
    a[5000] = 13;
}
```

```
80483b7: c7 05 60 e3 04 0d movl $0xd,0x804e360
```



- Sends `SIGSEGV` signal to user process
- User process exits with "segmentation fault"

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

13

251

Linking

- Linking
- Case study: Library interpositioning

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

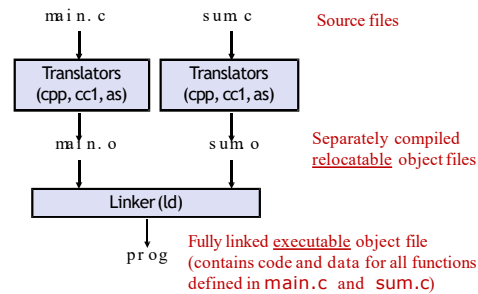
1

252

Static Linking

- Programs are translated and linked using a compiler driver:

```
linux> gcc -Og -o prog main.c sum.c
linux> ./prog
```



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



253

Why Linkers?

- Reason 1: Modularity**
 - Program can be written as a collection of smaller source files, rather than one monolithic mass.
 - Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library
- Reason 2: Efficiency**
 - Time: Separate compilation**
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
 - Space: Libraries**
 - Common functions can be aggregated into a single file...
 - Yet executable files and running memory images contain only code for the functions they actually use.



254

What Do Linkers Do?

- Step 1: Symbol resolution**
 - Programs define and reference symbols (global variables and functions):


```
void swap() {...} /* define symbol swap */
swap(); /* reference symbol swap */
int *xp = &x; /* define symbol xp, reference x */
```
 - Symbol definitions are stored in object file (by assembler) in symbol table.
 - Symbol table is an array of structs
 - Each entry includes name, size, and location of symbol.
 - During symbol resolution step, the linker associates each symbol reference with exactly one symbol definition.**

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



255

What Do Linkers Do? (cont)

- Step 2: Relocation**
 - Merges separate code and data sections into single sections
 - Relocates symbols from their relative locations in the .o files to their final absolute memory locations in the executable.
 - Updates all references to these symbols to reflect their new positions.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



256

Three Kinds of Object Files (Modules)

- Relocatable object file (.o file)
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each .o file is produced from exactly one source (.c) file
- Executable object file (.out file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
- Shared object file (.so file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called Dynamic Link Libraries (DLLs) by Windows

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



257

Linker Symbols

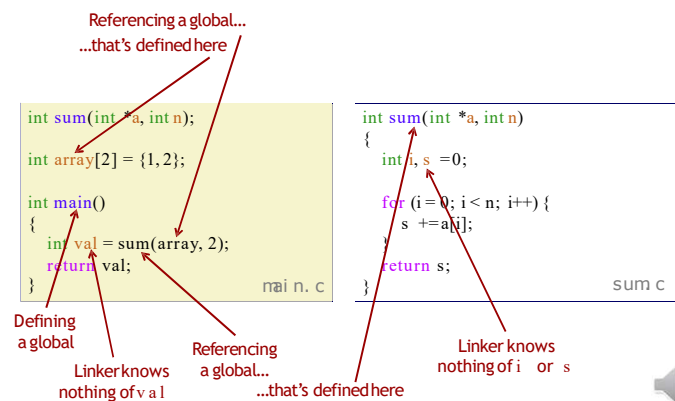
- Global symbols
 - Symbols defined by module m that can be referenced by other modules.
 - E.g.: non-static C functions and non-static global variables.
- External symbols
 - Global symbols that are referenced by module m but defined by some other module.
- Local symbols
 - Symbols that are defined and referenced exclusively by module m.
 - E.g.: C functions and global variables defined with the static attribute.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



258

Step 1: Symbol Resolution



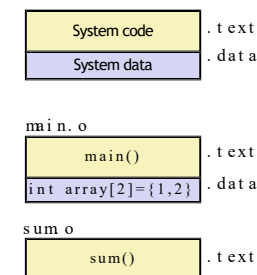
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




259

Step 2: Relocation

Relocatable Object Files



Executable Object File

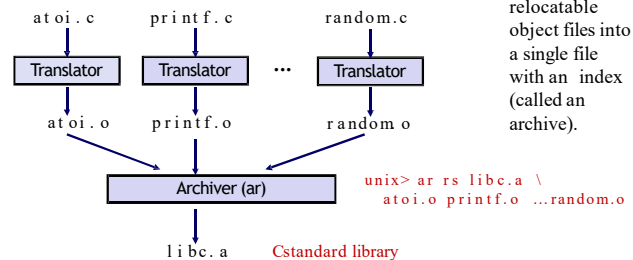


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



260

Creating Static Libraries



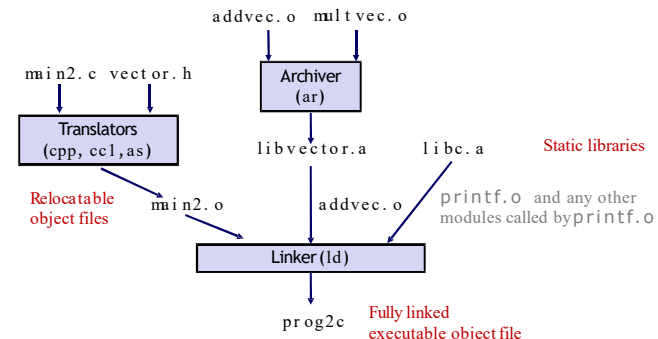
- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

25

261

Linking with Static Libraries



"c" for "compile-time"

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

262

Using Static Libraries

- Linker's algorithm for resolving external references:
 - Scan .o files and .a files in the command line order.
 - During the scan, keep a list of the current unresolved references.
 - As each new .o or .a file, obj, is encountered, try to resolve each unresolved reference in the list against the symbols defined in obj.
 - If any entries in the unresolved list at end of scan, then error.

- Problem:
 - Command line order matters!
 - Moral: put libraries at the end of the command line.

```

    unix> gcc -L. libtest.o -lmine
    unix> gcc -L. -lmine libtest.o
    libtest.o: In function 'main':
    libtest.o(.text+0x4): undefined reference to 'libfun'
  
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

29

263

Modern Solution: Shared Libraries

- Static libraries have the following disadvantages:
 - Duplication in the stored executables (every function needs libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink
- Modern solution: Shared Libraries
 - Object files that contain code and data that are loaded and linked into an application dynamically, at either load-time or run-time
 - Also called: dynamic link libraries, DLLs, .so files

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

264

Shared Libraries (cont.)

- Dynamic linking can occur when executable is first loaded and run (load-time linking).
 - Common case for Linux, handled automatically by the dynamic linker (`ld-linux.so`).
 - Standard C library (`libc.so`) usually dynamically linked.
- Dynamic linking can also occur after program has begun (run-time linking).
 - In Linux, this is done by calls to the `dlopen()` interface.
 - Distributing software.
 - High-performance web servers.
 - Runtime library interpositioning.
- Shared library routines can be shared by multiple processes.
 - More on this when we learn about virtual memory

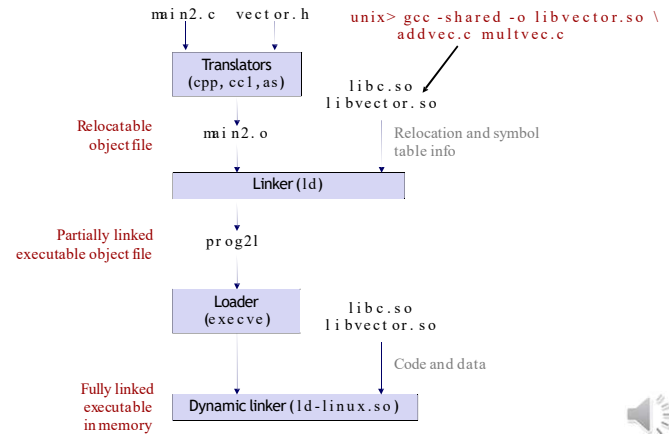
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



31

265

Dynamic Linking at Load-time



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



32

266

Linking Summary

- Linking is a technique that allows programs to be constructed from multiple object files.
- Linking can happen at different times in a program's lifetime:
 - Compile time (when a program is compiled)
 - Load time (when a program is loaded into memory)
 - Run time (while a program is executing)
- Understanding linking can help you avoid nasty errors and make you a better programmer.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



35

267